



Adopting Prometheus the Hard Way

Tim Simmons - Engineer - DigitalOcean
@timsimlol



Who am I

Tim Simmons

Engineer/Prometheus person

Observability Platforms

DigitalOcean

@timsimlol



Tim Simmons @timsimlol · Nov 13

what would win us president george washington or a single crunchy limon flamin hot cheeto



1



3





#goals

Understand what Prometheus **is**

Understand the **value** of a healthy Observability culture

How Prometheus (and friends) can **help**

Learn the **nuances** of scaling Prometheus and its adoption

Make Prometheus **successful** for you



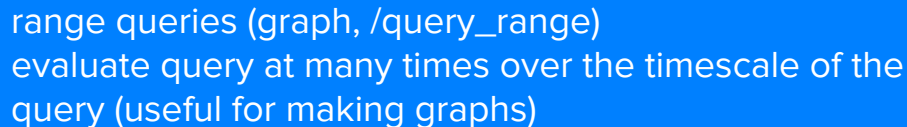
Prometheus

timeseries data (metrics)

```
timeseries = {(t0, v0), (t1, v1) .... }  
temperature = {(00:00, 20), (00:30, 22), ... (12:00, 25)}
```

queries (PromQL)

```
temperature = 25  
temperature[1h] = [(11:30, 23), (12:00, 25)]  
avg_over_time(temperature[12h]) = 20.94734
```





Prometheus - Labels

query	results
temperature	temperature{city="seattle", state="WA"} = 11 temperature{city="tacoma", state="WA"} = 10 temperature{city="san francisco", state="CA"} = 16 temperature{city="san jose", state="CA"} = 12 temperature{city="austin", state="TX"} = 15
temperature{state="CA"}	temperature{city="san francisco", state="CA"} = 16 temperature{city="san jose", state="CA"} = 12
avg(temperature) by (state)	temperature{state="WA"} = 10.5 temperature{state="CA"} = 14 temperature{state="TX"} = 15



Prometheus - Metrics

Pull based

Applications expose metrics HTTP endpoint

```
← → ↻ ⓘ localhost:9100/metrics
# HELP node_network_receive_bytes_total Network device statistic receive_bytes.
# TYPE node_network_receive_bytes_total counter
node_network_receive_bytes_total{device="br-58cdd73300bb"} 0
node_network_receive_bytes_total{device="br-c8e264c57cfe"} 0
node_network_receive_bytes_total{device="docker0"} 2.5736e+06
node_network_receive_bytes_total{device="enp3s0"} 0
node_network_receive_bytes_total{device="lo"} 1.4661817e+07
node_network_receive_bytes_total{device="tun0"} 1.8839546e+07
```



Prometheus - Exporters

JMX
Consul
ElasticSearch
Memcached
MongoDB
MSSQL
MySQL
PostgreSQL
ProxySQL
Redis
node/system metrics
NVIDIA
Ubiquiti
Kubernetes

Kafka
MQTT
RabbitMQ
Ceph
Gluster
Hadoop
Apache
HAProxy
Nginx
Varnish
Cloudflare
DigitalOcean
Docker
Fluentd

Go
Java or Scala
Python
Ruby
Bash
C++
Common Lisp
Elixir
Erlang
Haskell
Lua for Nginx/Tarantool
.NET / C#
Node.js
Perl
PHP
Rust



Prometheus - Cool Stuff

Alerts

if temperature > 30 for 3 days, send me an alert!

Relabeling

`temperature{location="Seattle, WA"} => temperature{city="seattle", state="WA"}`

Recording Rules

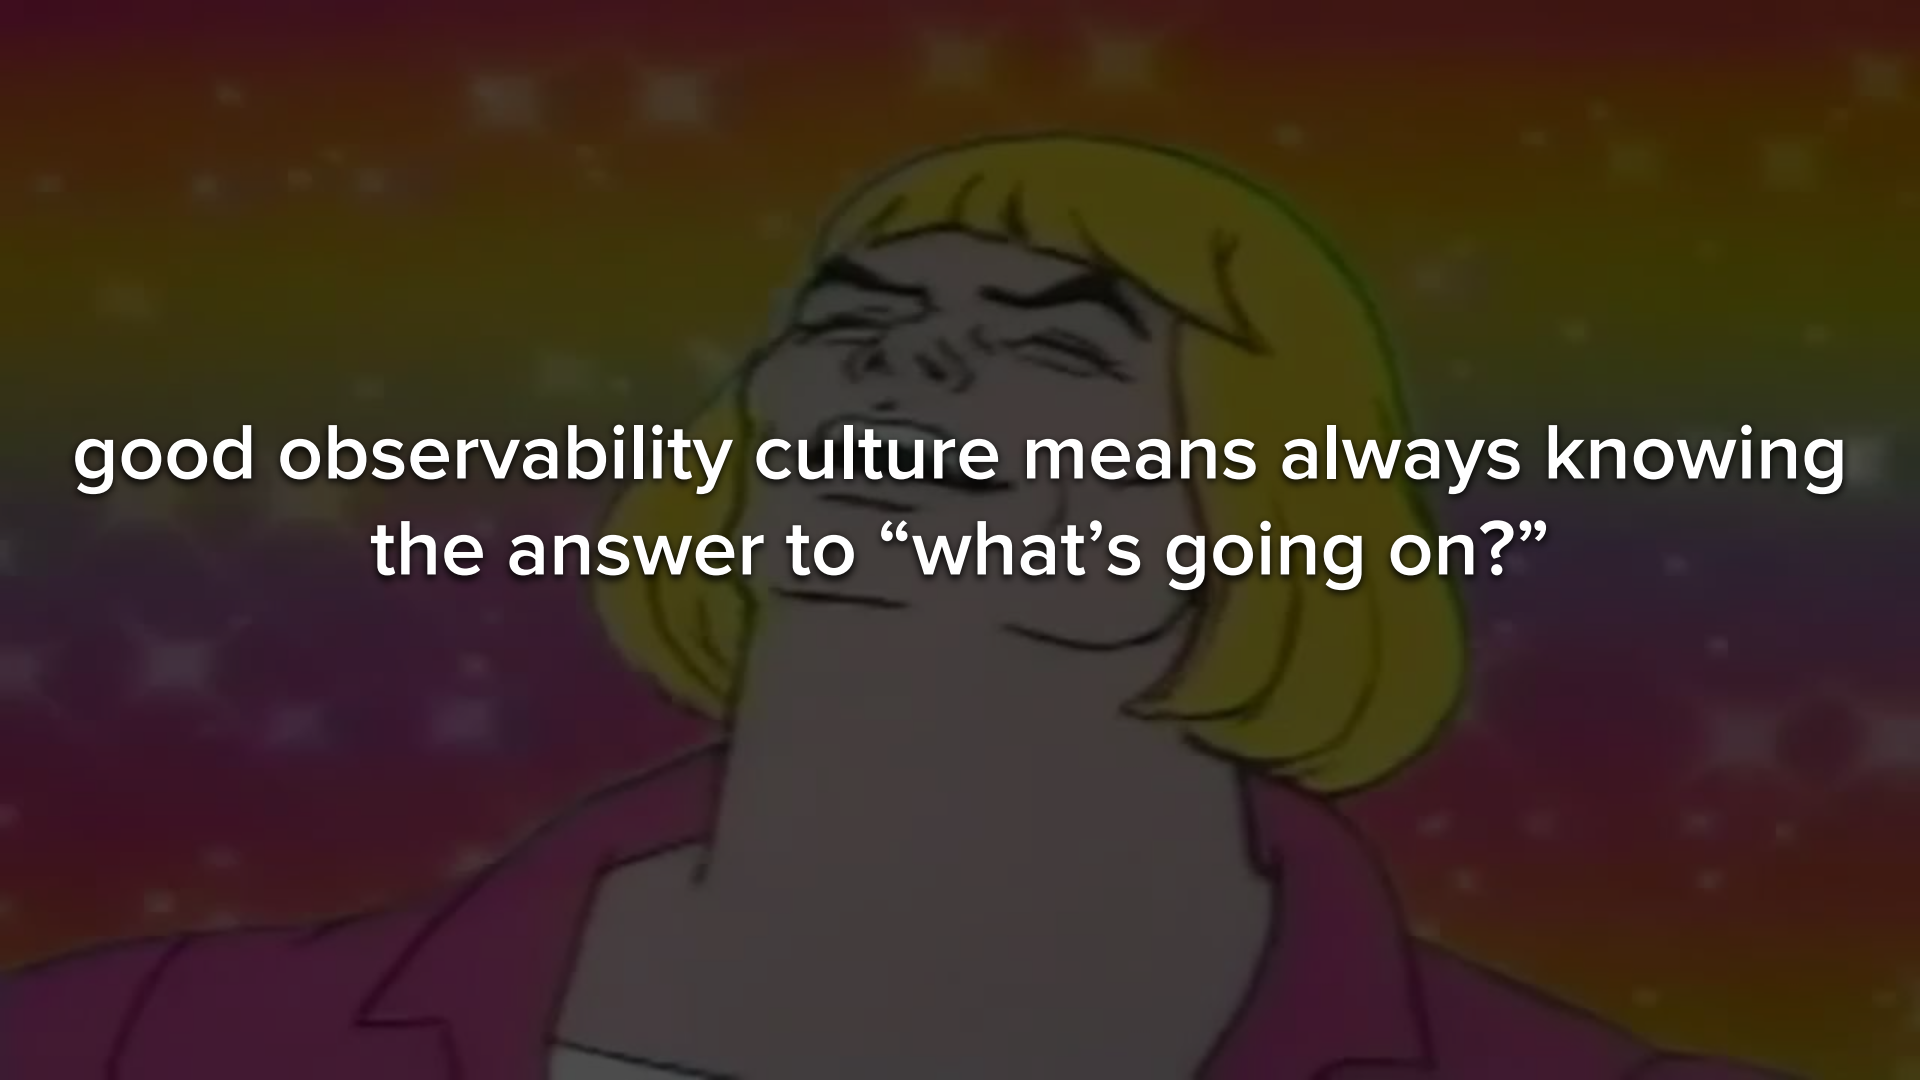
compute expensive queries regularly and save results to new metrics

Robust, Extendable Service Discovery

Kubernetes, GCE, Azure, OpenStack, EC2, Consul, DNS, Custom, etc

Endless Customization

federation, alerting routes/receivers/inhibits/integrations, scraping, remote read/write, limits



good observability culture means always knowing
the answer to “what’s going on?”



Measure everything, and then use that
data to make your life better





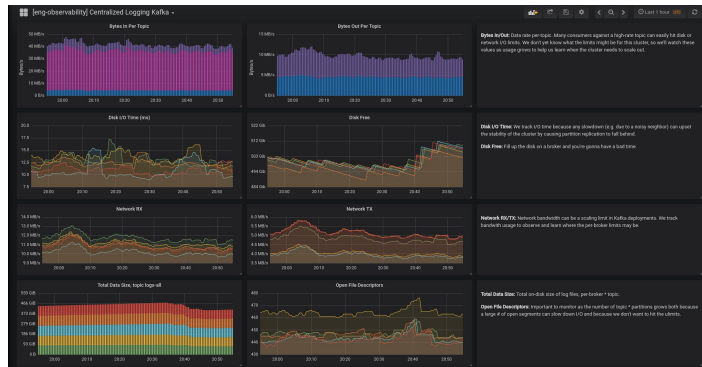
Dashboards are good

In an incident, “what’s going on” stands out 

Democratizing knowledge of “what’s going on” 

Onboarding 

More than just graphs!  





Friends of Prometheus - Grafana



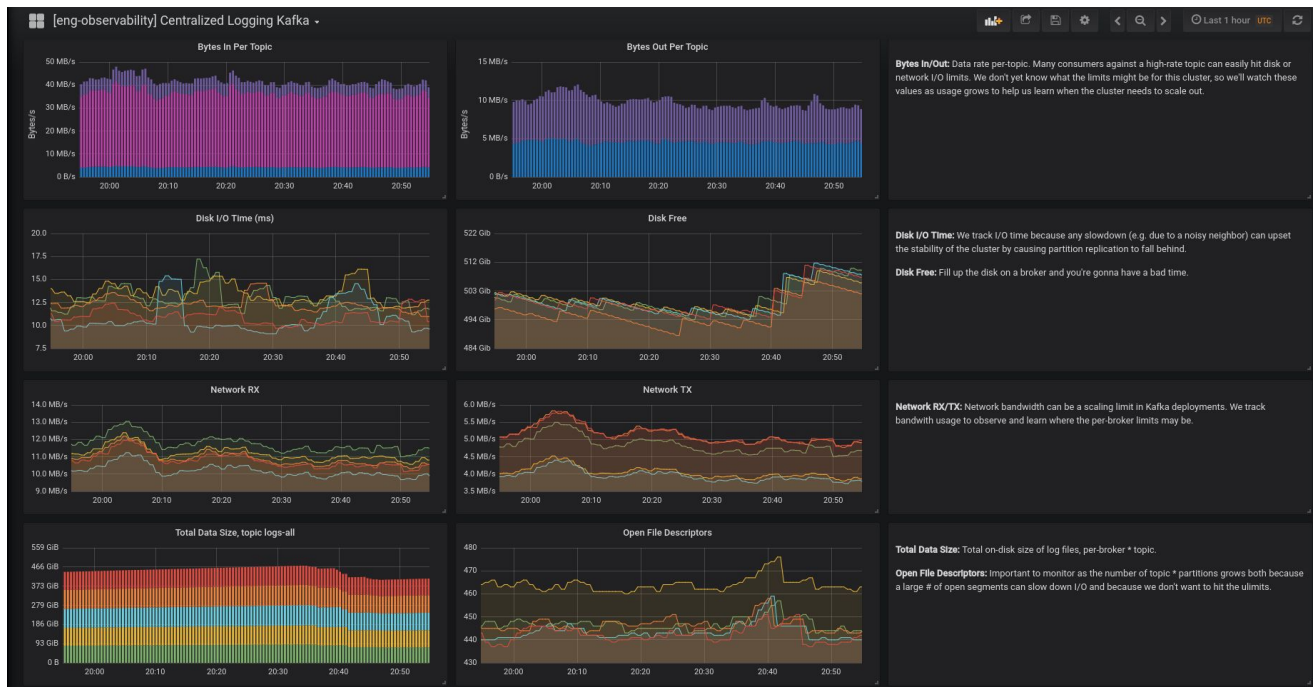
Dashboards

JS ✨

login/teams/rbac

API

More than just
Prometheus!



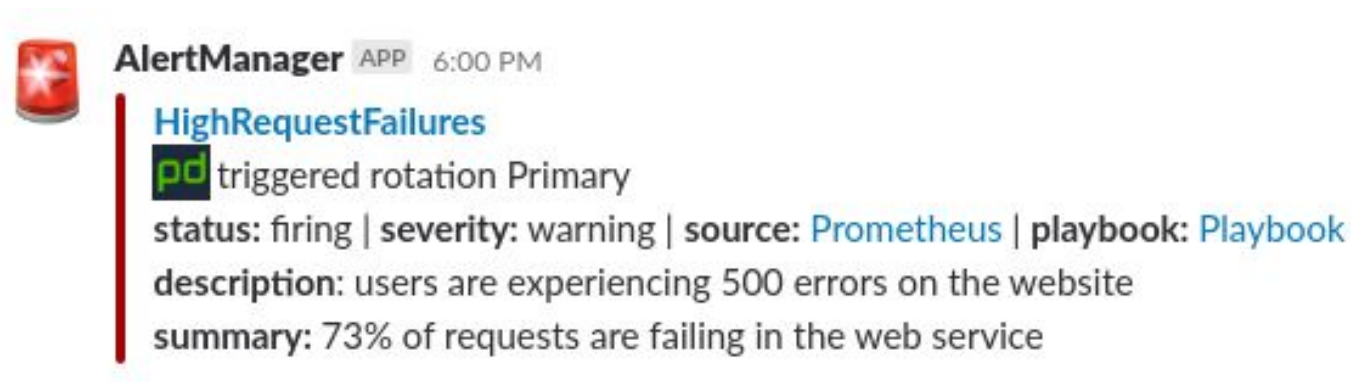


Alerts are good

You know when things break 🤔

A good alert can be handled by anyone 📱

More alerts -> more context 🧠





Friends of Prometheus - Alertmanager

Alertmanager is an **awesome** friend


Alerts are PromQL queries

Prometheus evaluates alerts on a cadence

If alert is true, send to Alertmanager

Deduplicate, Group, Notify

HA capabilities mean you never miss an alert





Observability Culture - Software

Build with Observability in mind rather than adding it later

Increasing returns as you build new systems

Clearer code organization/architectures

Quantify customer experience

SLOs and Error Budgets are right there

Make data driven decisions





@



DigitalOcean

STORY TIME!





Stories - Service Specific Dashboards/Alerts

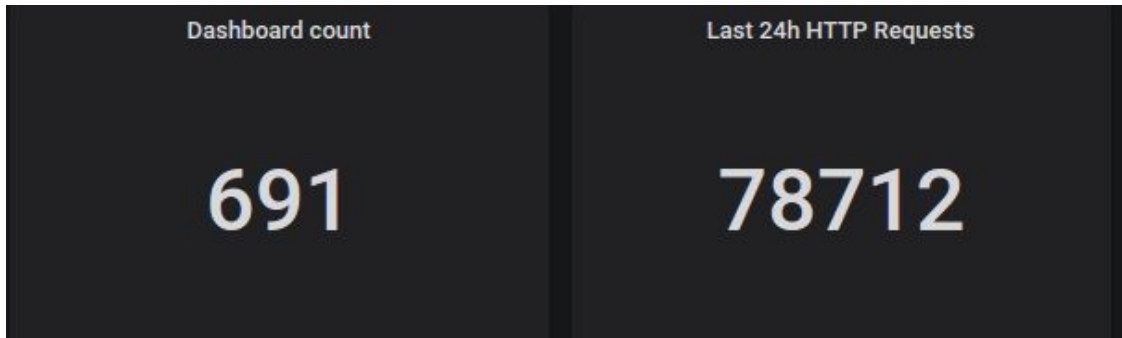
Grafana

30+ Team Directories

700+ Dashboards

Alerts!

Searchability 100



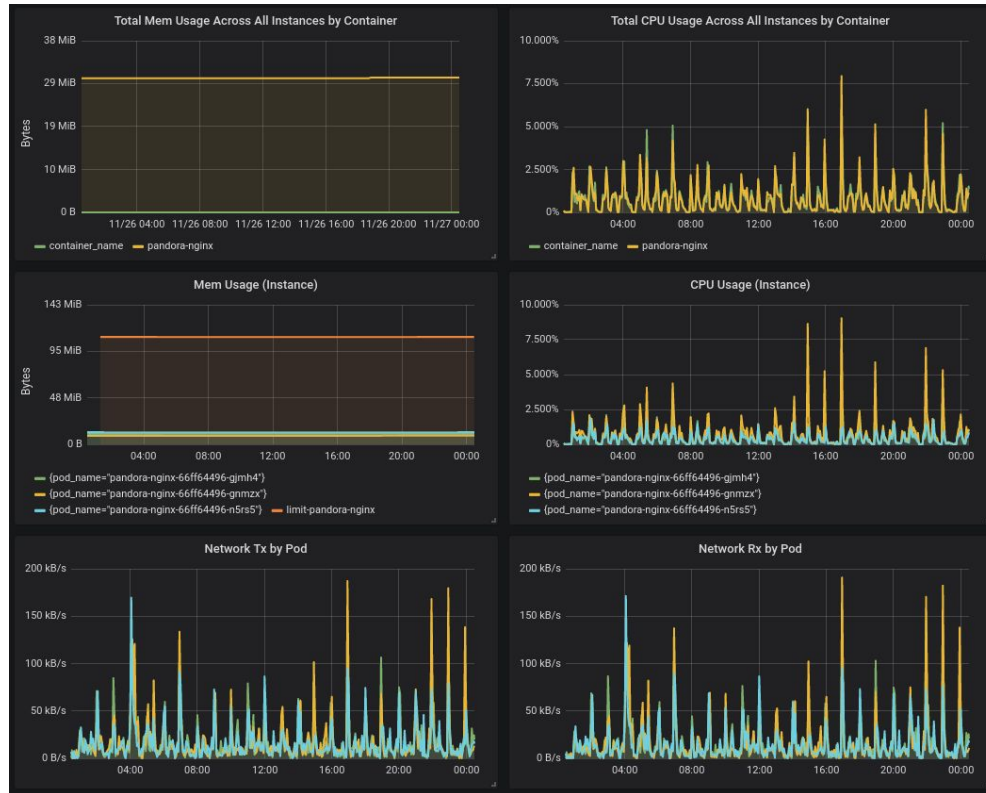


Stories - Shared Dashboards

Grafana templating allows systems with similar metrics to share a dashboard







k8s applications CPU/RAM/Net

Databases



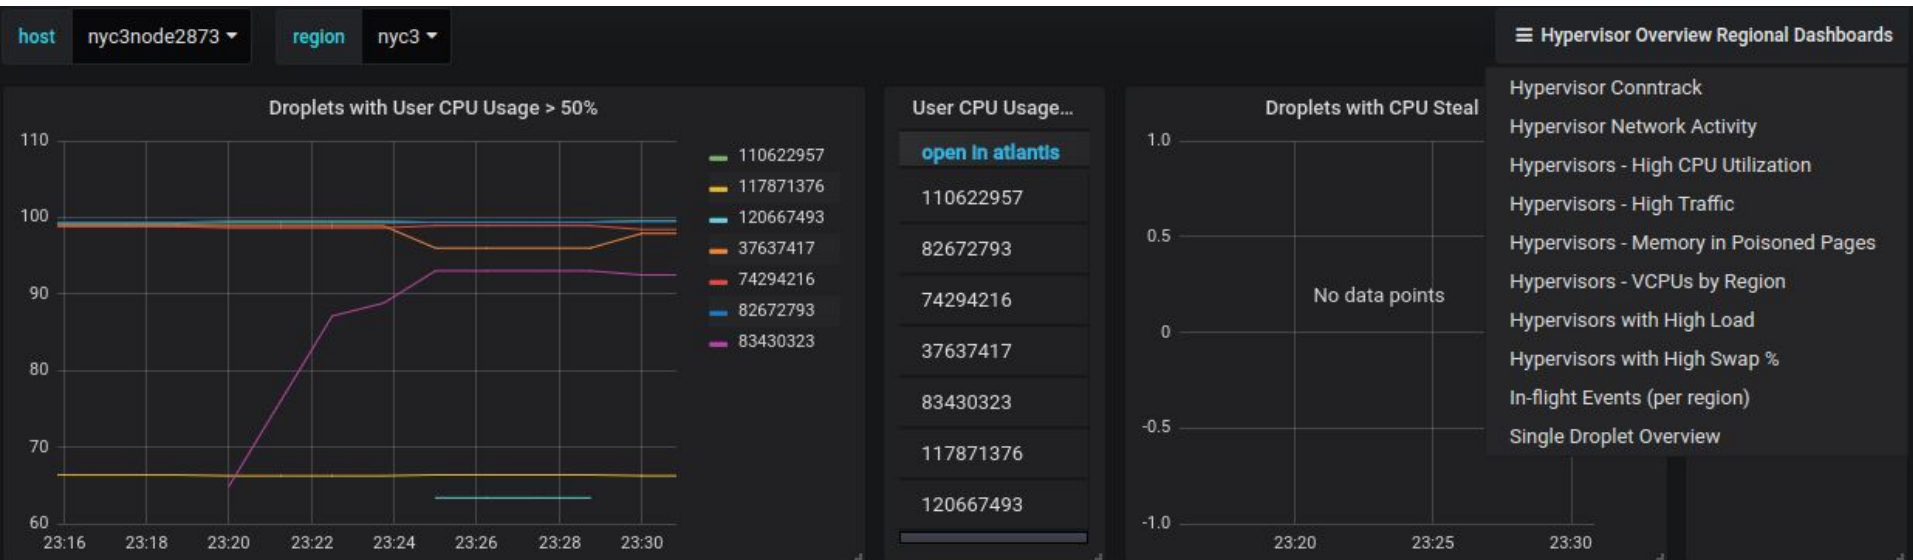


Stories - Hypervisor Metrics





Stories - Fleet Overview Graphs



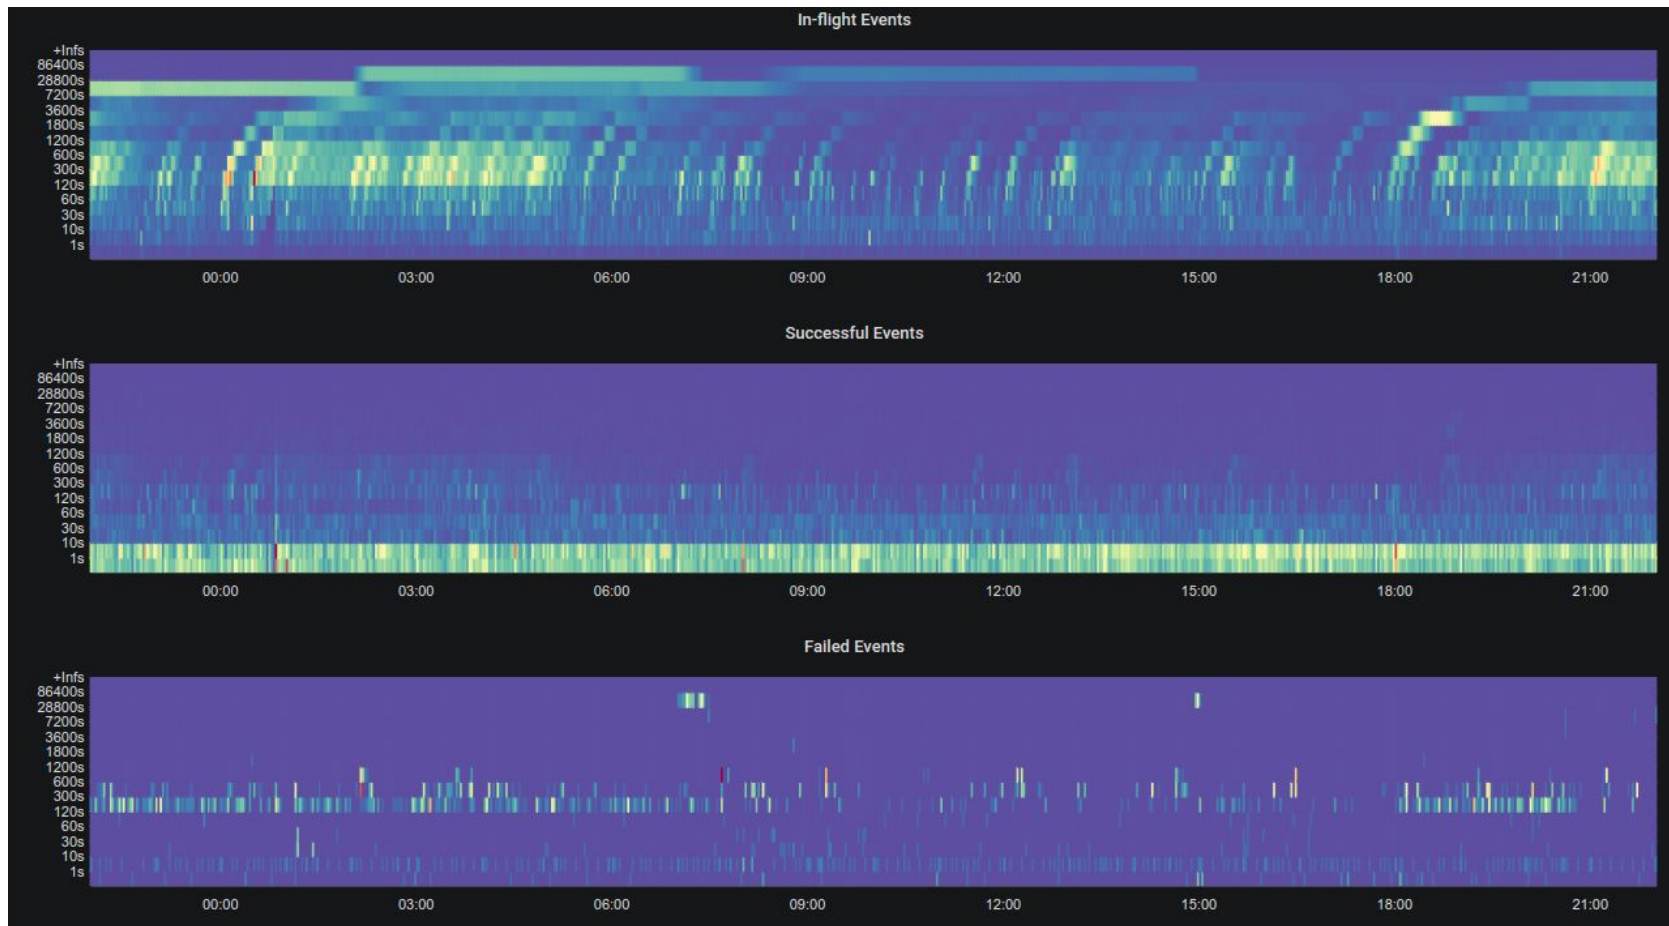


Stories - Droplet Metrics





Stories - Event Visualization





Stories - SLM/SLOs

Droplet

Compute SLM Month-to-Date

25m ago

Run




FILTERS

Monthly Charts is in the past 4 months

Weekly/Hourly Charts is in the past 5 weeks

hydropletmetrics

Aggregated Monthly Uptime by Region



Droplet Creates


Success Rates + Duration

Droplet Create Success Rate - monthly

Total Droplet Creates



Droplet Create Success Rate - by region, monthly






Stories - Capacity

Custom exporters


Relabel configs

Alerts!

Standard Capacity per region



High-CPU Capacity per region





Prometheus at DigitalOcean

192 Prometheus servers

200M+ time series

2M+ samples/second



Scaling Prometheus





If you take one thing away from this talk

Every **permutation of labels** in Prometheus creates a new time series


Individual queries should use **hundreds not thousands** of time series (at most)

Queries that **operate on** thousands of time series will overload Prometheus

Work out your query in the **Console** before graphing

Avoid **high cardinality** labels*

*unless you *really* know what you're doing



```
query: requests_total{path=~"/status|/)", method=~"(GET|POST)"}
{__name__="requests_total", path="/status", method="GET", instance="10.0.0.1:80"}
{__name__="requests_total", path="/status", method="POST", instance="10.0.0.3:80"}
{__name__="requests_total", path="/", method="GET", instance="10.0.0.2:80"}
```




Leave Headroom

Prometheus needs **memory** headroom to execute queries

High cardinality metrics, big dashboards, long retention, concurrent queries

Physical Resources (CPU/Memory/Disk) ⚙️ 🗑️





Shard Metrics Functionally

Metrics will get too big to query

10s of millions of timeseries on a single server



You will need to shard the metrics across multiple Prometheus servers

Pick a dimension that is a query boundary

Never split metrics that you want to query together

Split on region, service, team

Not instance, application



HA Prometheus Pairs


Deploy Prometheus in pairs (at least?)

Scrape the same metrics

Proxy queries *Active-Active*, *Active-Passive*

Proxies are good

If something bad happens, you're good(ish)!





Find Bad Metrics

<https://www.robustperception.io/which-are-my-biggest-metrics>

```
topk(10, count by ( name ) ({ name =~".+"}))
```

```
topk(10, count by ( name , job) ({ name =~".+"}))
```

```
topk(10, count by (job) ({ name =~".+"}))
```

```
sum(scrape samples scraped) by (job)
```

Expensive, but worth it



Friends of Prometheus - Trickster



github.com/Comcast/trickster

Caches query results

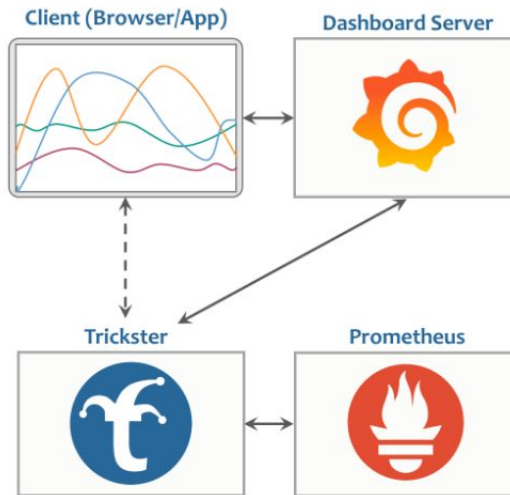
(Partial, or full)

Step boundary normalization

Trickster build passing

go report A+

Trickster is a reverse proxy cache for the [Prometheus HTTP API v1](#) that dramatically accelerates dashboard rendering times for any series queried from Prometheus.





Scaling Prometheus Adoption





Creating Value

Engineers (Dev/SRE/Support) are paid to create business value

New Features

Maintenance

Every engineer has customers





Maintenance

Noun

the process of maintaining or preserving something

“revenue protection”*

Measure performance

Debug poor performance

Support users

This is what **Observability** tools do!



*<https://twitter.com/DarkAndNerdy/status/1070690952472350720>



Revenue Protection

Systems are rarely built with “revenue protection” in mind

Chasing features to provide value is natural

New shiny can tarnish the old dull

Observability tools help protect the 💰





So if you want to provide all the value...

Build Observability into your application in the **beginning** under *pressure* to deliver the new shiny

OR

Retrofit Observability **later**, under *pressure* because something is broken or hard to maintain





Adding Observability with Prometheus

With limited time be efficient as possible, Ctrl+c, Ctrl+v!

In an organization, **patterns proliferate**

Prometheus metrics are custom for every application

Metrics data quality is **more** dependent on your decisions



Prometheus is deceptively simple

Basics are breezy

Creating meaningful custom metrics is harder

Operations are easy at first

Defining/Delivering quality alerts

Making great dashboards

Dealing with expensive queries

Be careful copying those patterns



If anyone gets this one come talk to me after so we can be best friends



On the Path to Maintenance

Every decision you make is a multiplier on your ability to efficiently maintain a system.



But it is OK

Prometheus is great!

You will get results, no matter what.

Done well, you can 10000x those results tho

Especially in medium+ size orgs



You need people for this.



You need people to do this

What

Ensure engineers can *effectively utilize* limited Observability time

Accelerate *basic* Prometheus understanding

How

Operate

Centralize Knowledge and Configuration

Consult



Operate

“own” Prometheus and friends

Make sure no one else has to operate Prometheus

Upgrade and improve

Be on-call

Watch out for metric/alerting regressions

Make sharding decisions






Centralize Configuration

Make it easy to copy good patterns

Chef, Ansible, Git, Docker

Provide a **good** abstraction





Centralize Knowledge

Standardize and document what to measure (Golden Signals, RED, USE)

Establish patterns around similar use cases (RPC, Physical Resources, k8s)

Document good examples (client usage, dashboards, alerts)

Document how to alert good (“On-call doesn’t have to suck” -Cindy Sridharan @copyconstruct
<https://medium.com/@copyconstruct/on-call-b0bd8c5ea4e0>)

Document how to use your internal tools



What I Did

- YAML abstractions
- GitOps
- Observability reviews PRs
- Created confusion and delay



```
---
hosts:
  prod-elauneind-nova:
    port: 9123
    kubernetes:
      cluster: nova
      namespace: observability
      application: elauneind

alerts:
- alert: Elauneind Prod Down
  expr: |-
    sum(up{job="prod-elauneind"}) == 0
  for: 10m
  labels:
    team: eng-observability
    service: elauneind
    severity: warning
  annotations:
    description: "elauneind down in prod {{ $labels.cluster }}"
```

```
---
owner: eng-observability@digitalocean.com
slack_channels:
  - "#observability-alerts"
pagerduty_keys:
  - momsspaghetti
severity_routing:
  warning:
    - "#observability-alerts"
  critical:
    - "#observability-alerts"
    - momsspaghetti
```



What I Did

pandora-users

The endpoint for adding/modifying services to be scraped by Pandora, the hosted metrics service from the Observability team ❤️

Fork this repository and make a PR to get started!

- [Scrape my metrics right now!](#)
- [Send me alerts right now!](#)
- [Send CloudOps alerts right now!](#)
- [What is Pandora anyway?](#)
 - [What is Prometheus anyway?](#)
- [Key Concepts](#)
 - [Teams](#)
 - [Services](#)
 - [A Note On Service Names](#)
 - [Groups](#)
- [Validating Inputs](#)
 - [Pandora CLI](#)
 - [Validating Pandora Users YAML](#)
 - [Generate Prometheus config files](#)
 - [promtool](#)
 - [Alerting Routing Tree](#)
 - [Validation TODOs](#)
- [Advanced Concepts](#)
 - [Alerts](#)
 - [Alerts Best Practices](#)
 - [Simple Alert Delivery](#)
 - [Receivers](#)
 - [Slack](#)
 - [Pagerduty](#)
 - [Webhooks](#)
 - [Routes](#)
 - [Inhibit Rules](#)
 - [Recording Rules](#)
 - [Relabel Configurations](#)
 - [Custom Scraping Config](#)
 - [Host Group Discovery Methods](#)
 - [Static](#)
 - [Chef](#)
 - [Marathon](#)
 - [sds](#)
 - [Kubernetes](#)
 - [Consul](#)
- [How to Review and Promote a pandora-users PR](#)

Wrote a lot of
documentation

Pandora

Created by bknox, last modified by tsimmons on Aug 28, 2018

Pages You Should Read

[How to Query Prometheus with PromQL](#)

[Instrumenting Your Application with Prometheus](#)

[What Makes a Good Alert](#)

[What Makes a Good Metric and What Should I Measure?](#)

Other Pages You Might Like

[How Does Alertmanager Work in Pandora?](#)

[How to use Pandora to route RedAlert alerts from Centralized Logging](#)

[How to use the Prometheus API with Pandora](#)

[Pandora Architecture](#)

[Systems to Deprecate](#)

[Using Grafana with Pandora](#)

[Using node_exporter textfiles to export arbitrary metrics from a node](#)



What I Would Do Today

The screenshot shows the GitHub repository page for `coreos/prometheus-operator`. At the top, the repository name is displayed with navigation links for Watch (120), Star (1,988), and Fork (956). Below this is a navigation bar with links for Code, Issues (213), Pull requests (39), Projects (0), Wiki, and Insights. The repository description states: "Prometheus Operator creates/configures/manages Prometheus clusters atop Kubernetes" with a link to <https://coreos.com/operators/prometheus>. Below the description are tags for `kubernetes`, `prometheus`, and `monitoring`. At the bottom, statistics are shown: 2,481 commits, 25 branches, 48 releases, 246 contributors, and Apache-2.0 license.

Build administrative app on top of prometheus-operator

Configure scraping by pushing info from local app manifests (w/ CLI) into a stateless application that verifies/shards/places metrics by creating operator compatible CRDs



Consult

- Point out getting started resources
- Coach initial attempts
- Enforce good patterns
- Suggest advanced configuration
- Lend a hand in an incident
- Help solve larger organizational issues





You need people to do this

Operate

Centralize Knowledge

Centralize Configuration

Consult

Anyone can do this!

Manage Centralized Logging? Tracing? Exceptions?



What does adoption look like?

Some people will get it, love it, and get involved

Most will cypasta patterns

Some people will want the moon

- 1s sampling
- infinite retention
- massive label cardinality





Friends of Prometheus - LTS - Thanos/M3




Thanos

Prometheus sidecars

Backed by block storage

Query Prometheus and block storage

Downsampling! Global query view!



M3

Prometheus remote read/write


Replicated, distributed storage of metric data

Configurable time resolutions, metric durability

Support for multiple clusters based on etcd

Downsampling!

Global view of metrics





Conclusion

Prometheus is good, and has friends that make it even better

Having a healthy Observability culture has enormous value

Using Prometheus to successfully power an Observability culture is possible

You need people dedicated to making it good



The Greek Titan Prometheus stole fire, and gave it to humanity.

The computer thing Prometheus is trying to give you the ability to debug, monitor, and understand your systems.



Christian Griepenkerl (1839-1916): Die Strafe. Photo © Maicar Förlag - GML

The Greek Titan Prometheus was punished. He was chained to a rock and an eagle was sent to eat his regenerating liver every day for eternity.

Don't punish your Prometheus by letting it go unmanaged.

