



KubeCon

CloudNativeCon

— North America 2018 —

Protect Your Kubernetes Data

Friends Don't Let Friends Leave Their Kubernetes Data Unprotected

Rita Zhang
Principal Software Engineer @ Microsoft

@ritazzhang

Rita Zhang

- Software engineer @ Microsoft, San Francisco
- Kubernetes upstream features, Azure Kubernetes Service
- Maintainer for K8s KMS plugin for Azure Key Vault, Keyvault-flexvolume, aad-pod-identity



TRANSPORTATION

CARS

TESLA

Tesla's cloud was used by hackers to mine cryptocurrency

Mining bitcoin on Elon's dime

By Andrew J. Hawkins | @andyjayhawk | Feb 20, 2018, 1:39pm EST

The initial point of entry for the Tesla cloud breach, Tuesday's report said, was an unsecured administrative console for **Kubernetes**, an open source package used by companies to deploy and manage large numbers of cloud-based applications and resources.

A screenshot of a web browser displaying a Kubernetes Secrets page. The URL in the address bar is https://[REDACTED]/#/secret/default/aws-s3-credentials?namespace=default. The page title is "Name" and the logo is "kubernetes". A search bar is present. The main navigation menu shows "Config and storage > Secrets > aws-s3-credentials". On the left, a sidebar lists "Namespace" (default), "Overview", "Workloads" (including Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, Stateful Sets), "Discovery and Load Balancing" (Ingresses, Services), and "Config and Storage". The main content area is divided into "Details" and "Data". In the "Details" section, it shows Name: aws-s3-credentials, Namespace: default, Creation time: 2017-10-12T22:29, and Type: Opaque. In the "Data" section, two fields are shown: aws-s3-access-key-id: [REDACTED] and aws-s3-secret-access-key: [REDACTED]. The entire content area is framed by a thick black border.

RedLock

Name

Not Secure | https://[REDACTED]/#/secret/default/aws-s3-credentials?namespace=default

kubernetes

Search

Config and storage > Secrets > aws-s3-credentials

Namespace

default

Overview

Workloads

- Daemon Sets
- Deployments
- Jobs
- Pods
- Replica Sets
- Replication Controllers
- Stateful Sets

Discovery and Load Balancing

- Ingresses
- Services

Config and Storage

Details

Name: aws-s3-credentials

Namespace: default

Creation time: 2017-10-12T22:29

Type: Opaque

Data

aws-s3-access-key-id: [REDACTED]

aws-s3-secret-access-key: [REDACTED]

An attacker who can successfully access your cluster database can compromise your entire cluster and have access to your cloud resources.

Kubernetes Database

- Uses etcd as its persistent storage for API objects
- Stores secrets as base64 encoded plaintext

The security footgun in etcd



Giovanni Collazo

March 16, 2018

security



“Authentication was added in etcd 2.1. ... etcd before 2.1 was a completely open system; anyone with access to the API could change keys. In order to preserve backward compatibility and upgradability, this feature is off by default.” [Read more from coreos etcd doc](#)

<https://elweb.co/the-security-footgun-in-etcd/>

I did a [simple search on shodan](#) and came up with **2,284 etcd servers on the open internet.**

CREDENTIALS, a lot of CREDENTIALS. Credentials for things like cms_admin, mysql_root, postgres, etc. Passwords for databases of all kinds, AWS secret keys, and API keys and secrets for a bunch of services.

GET http://<ip address>:2379/v2/keys/?recursive=true

password	8781
aws_secret_access_key	650
secret_key	23
private_key	8

```
- {
  key: "/registry/secrets/ci",
  dir: true,
- nodes: [
- {
    key: "/registry/secrets/ci/docker-repo-key",
    value: "{\"kind\":\"Secret\",\"apiVersion\":\"v1\",\"metadata\":{\"name\":\"docker-repo-key\",\"namespace\":\"ci\",\"uid\":\"c98f370a-d825-11e6-a077-bc764e04\"},
modifiedIndex: 12910772,
createdIndex: 12910772
},
- {
    key: "/registry/secrets/ci/default-token-dt40q",
    value: "{\"kind\":\"Secret\",\"apiVersion\":\"v1\",\"metadata\":{\"name\":\"default-token-dt40q\",\"namespace\":\"ci\",\"uid\":\"f849clc5-ea76-11e7-94ad-bc76\"},
\"ca.crt\":\"LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUZVVENDQXptZ0F3SUJBZ0lKQU1ydVRyaThFTnBjTUEwR0NTcUdTSWIzRFFFQkN3VUFNRDh4Q3p8SkJnTlYK0
modifiedIndex: 651393897,
createdIndex: 651393897
}
],
modifiedIndex: 4005597,
createdIndex: 4005597
}
```



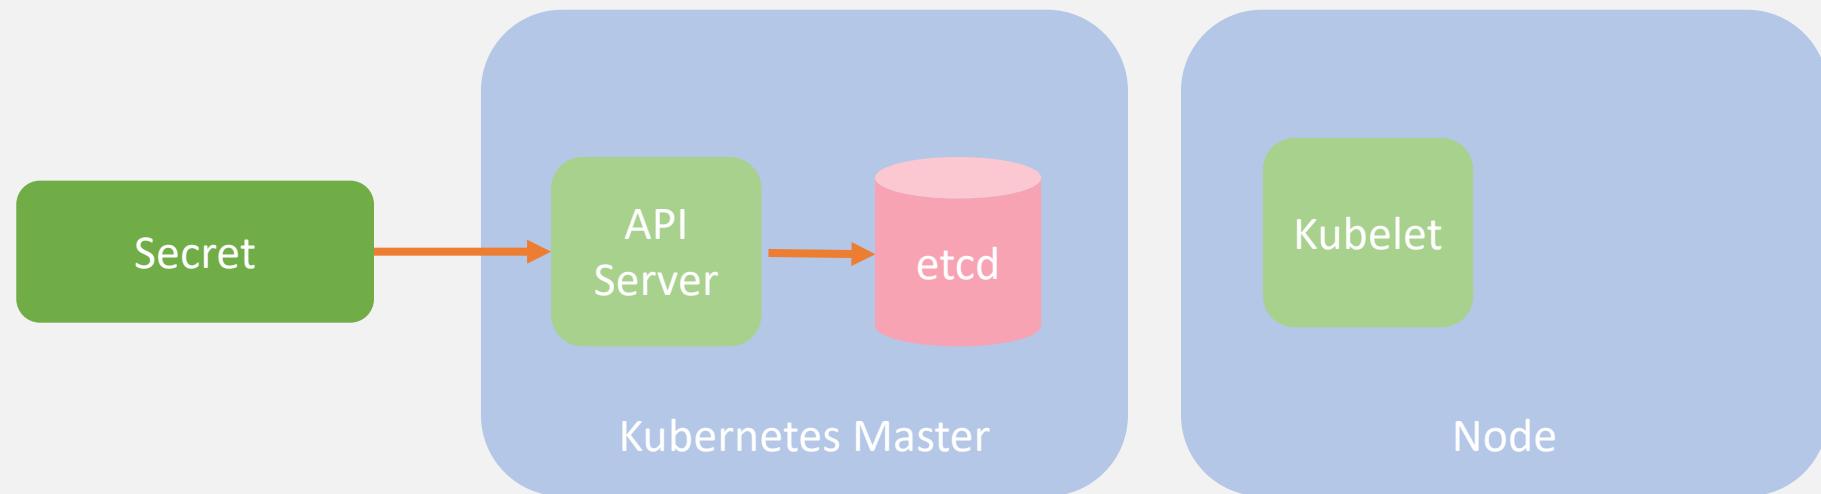
So...How do I secure my cluster?

- There are many things you can do
 - Control access to the Kubernetes APIs
 - Control access to the Kubelet
 - Control privileges containers run with
 - Restrict network access
 - Restrict resource access
 - Restrict access to etcd
 - **Encrypt etcd data at rest**
 - **Store application secrets outside of Kubernetes**
 - **Restrict access to resources with pod identity**

<https://kubernetes.io/docs/tasks/administer-cluster/securing-a-cluster/>

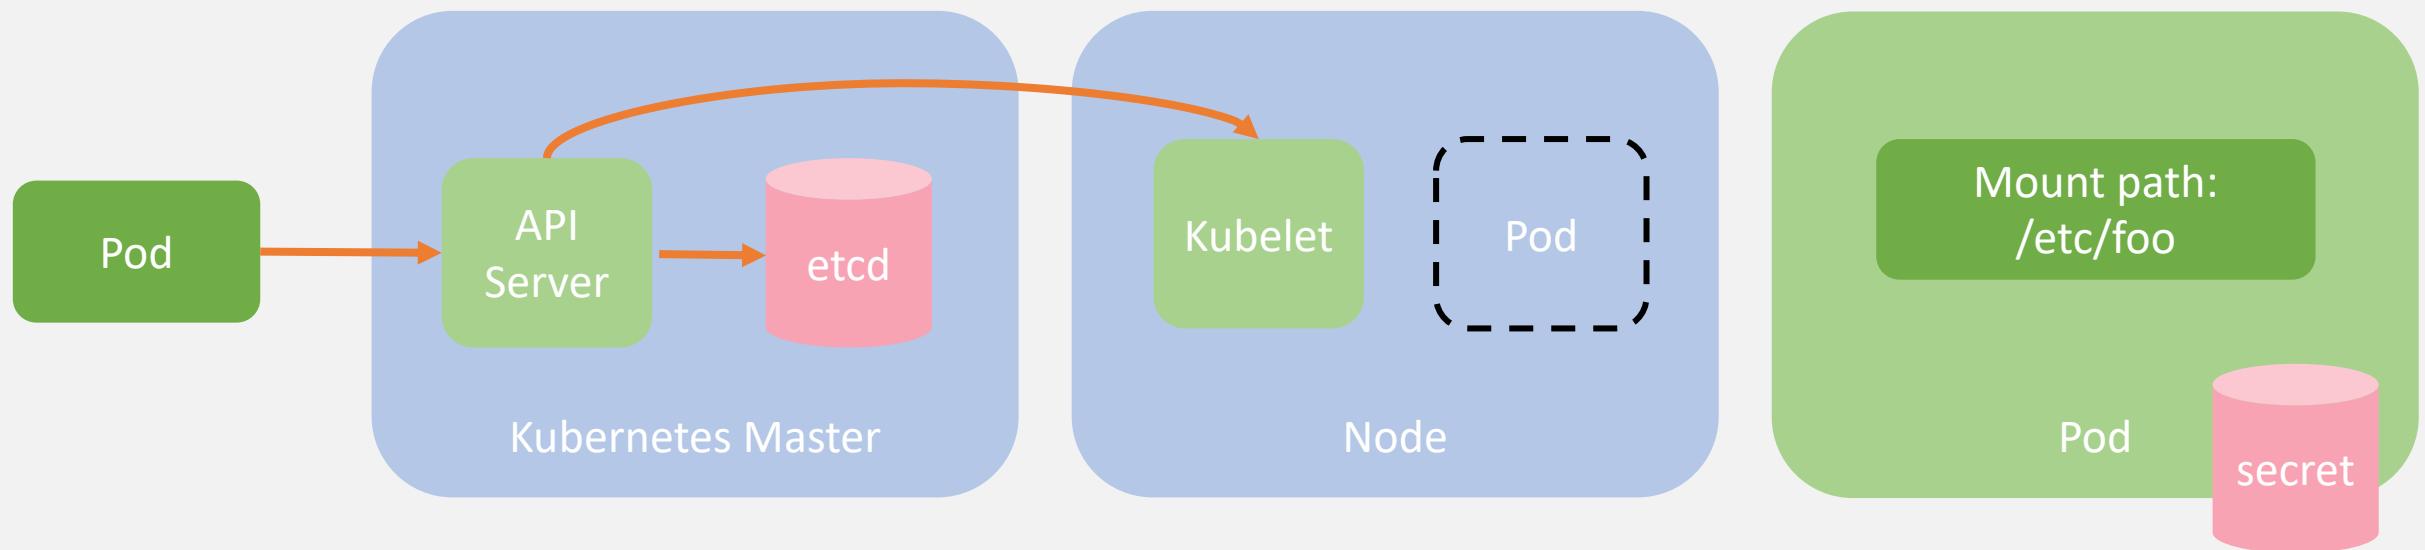
Secrets

```
kubectl create secret generic secret1
```

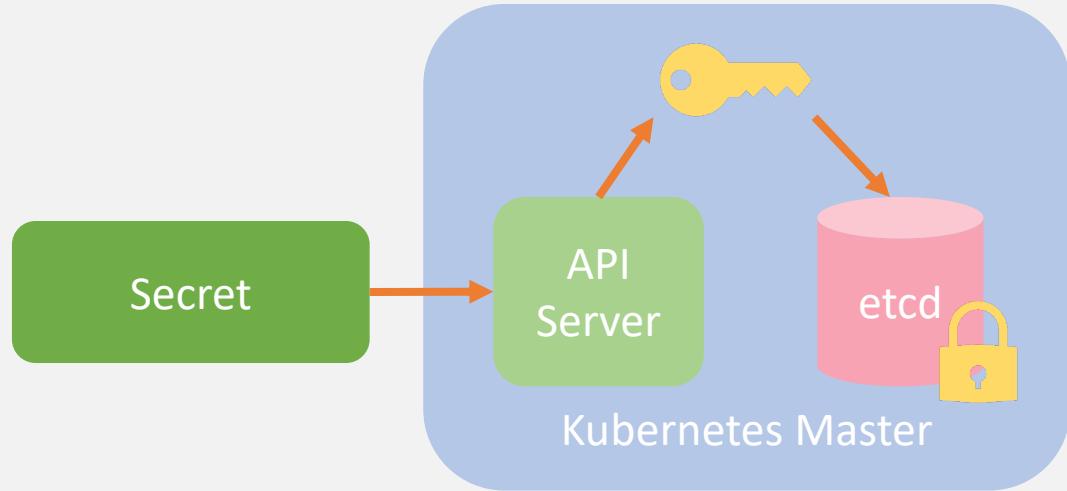


Pod using Secret

`kubectl create -f pod-using-secret.yaml`



Encryption at Rest

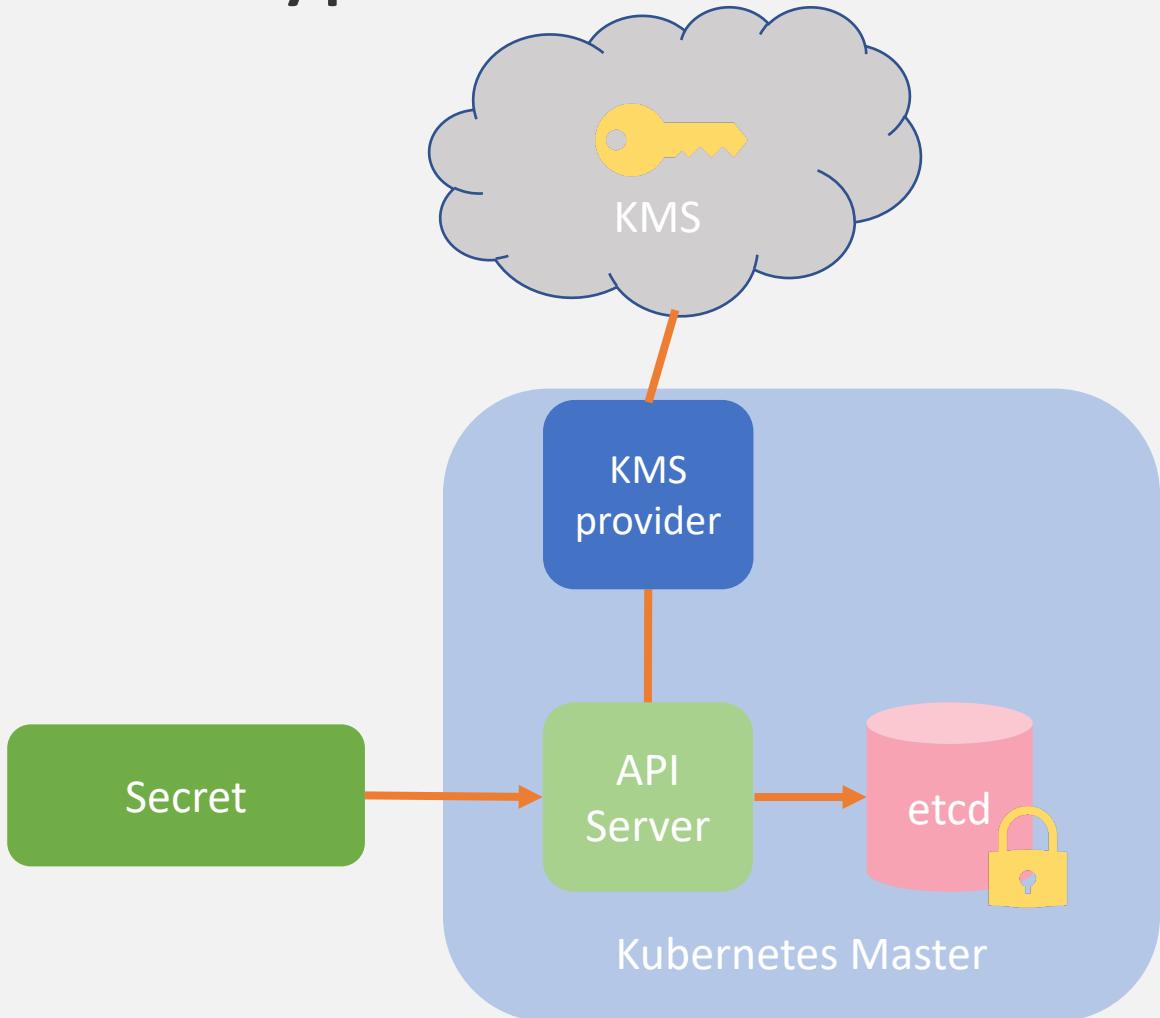


- Kubernetes v1.7+
- etcd v3 required
- encryption using keys in config file on master
- Plain text, encoded with base64

<https://kubernetes.io/docs/tasks/administer-cluster/encrypt-data/>

```
kind: EncryptionConfig
apiVersion: v1
resources:
  - resources:
    - secrets
providers:
  - aescbc:
    keys:
      - name: key1
        secret: <BASE 64 ENCODED SECRET>
  - identity: {}
```

Key Management Service (KMS) Provider for Encryption at Rest



- Kubernetes v1.10, v1.13 stable
- etcd v3 required
- Separate key management from K8s cluster management
- Supports encryption using keys stored in external trusted Key Management Service (KMS), e.g. Azure Key Vault, Google Cloud KMS
- HSM-protected keys

<https://kubernetes.io/docs/tasks/administer-cluster/kms-provider/>

Before V1.13

```
kind: EncryptionConfig
apiVersion: v1
resources:
- resources:
  - secrets
providers:
- kms:
    name: myKmsPlugin
    endpoint: unix:///tmp/socketfile.sock
    cachesize: 100
- identity: {}
```

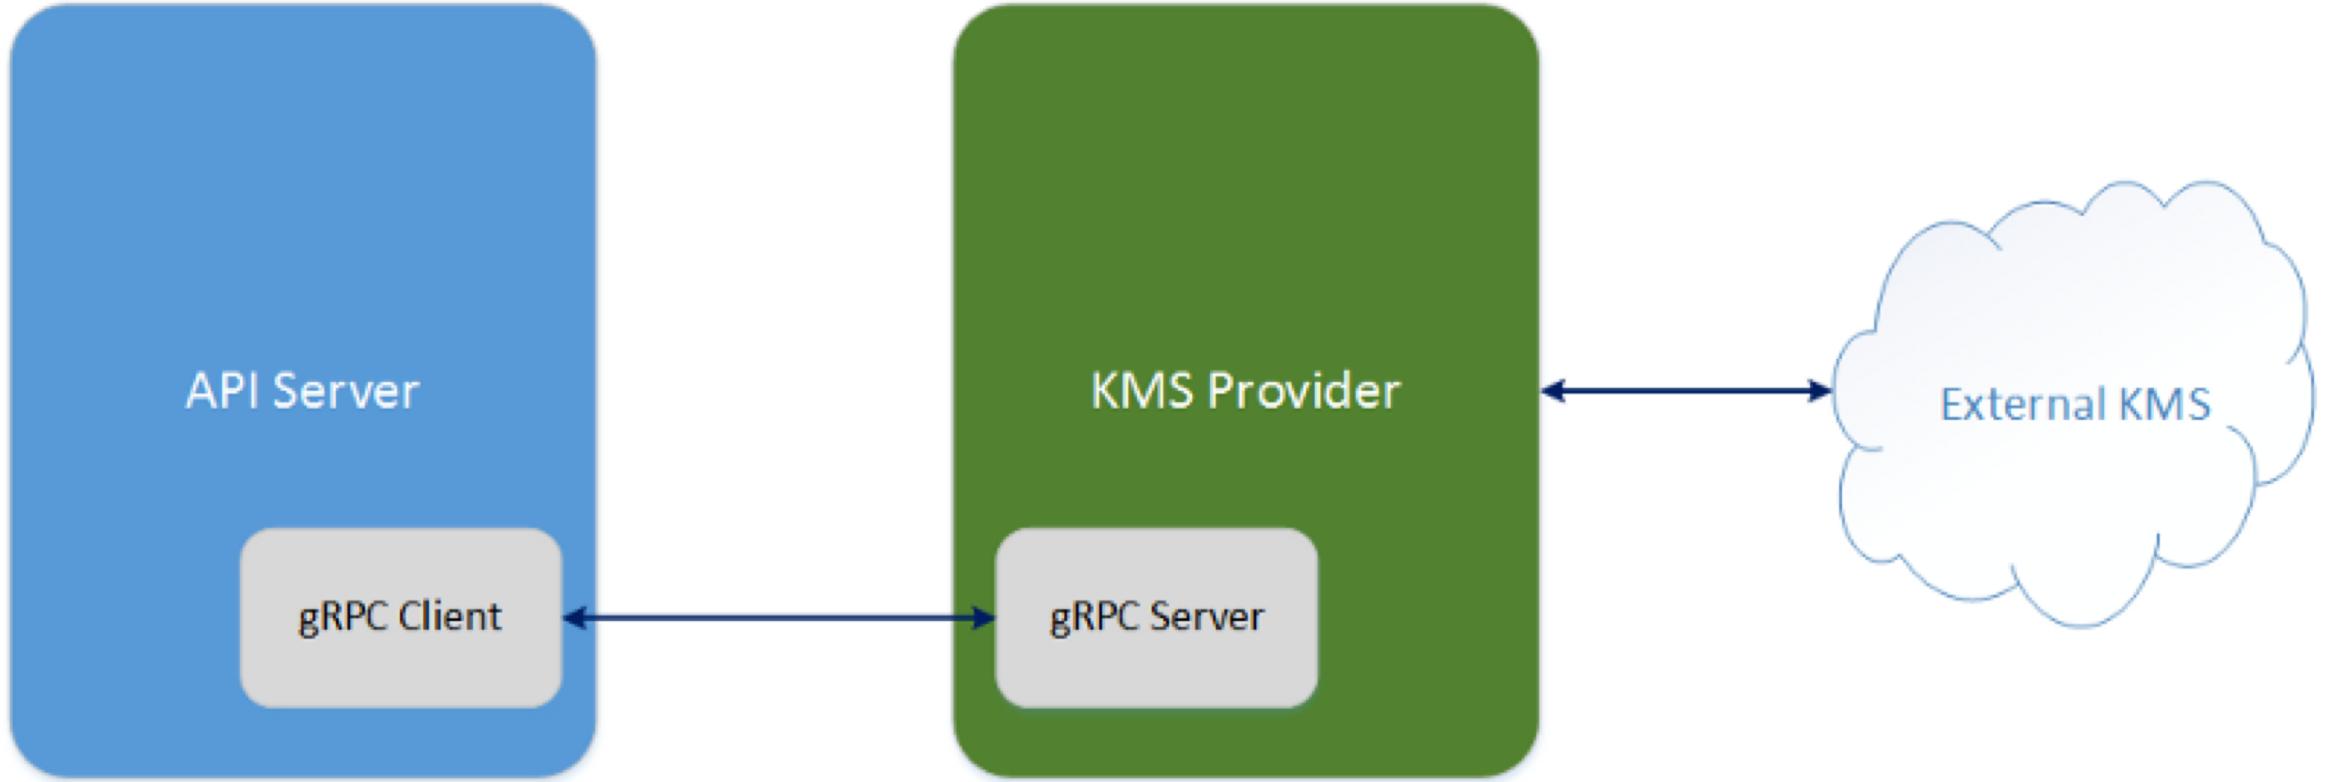
V1.13

```
kind: EncryptionConfiguration
apiVersion: apiserver.config.k8s.io/v1
resources:
  - resources:
    - secrets
  providers:
  - kms:
      name: myKmsPlugin
      endpoint: unix:///tmp/socketfile.sock
      cachesize: 100
  - identity: {}
```



Demo: K8s cluster with Azure Key Vault data encryption

High Level Design

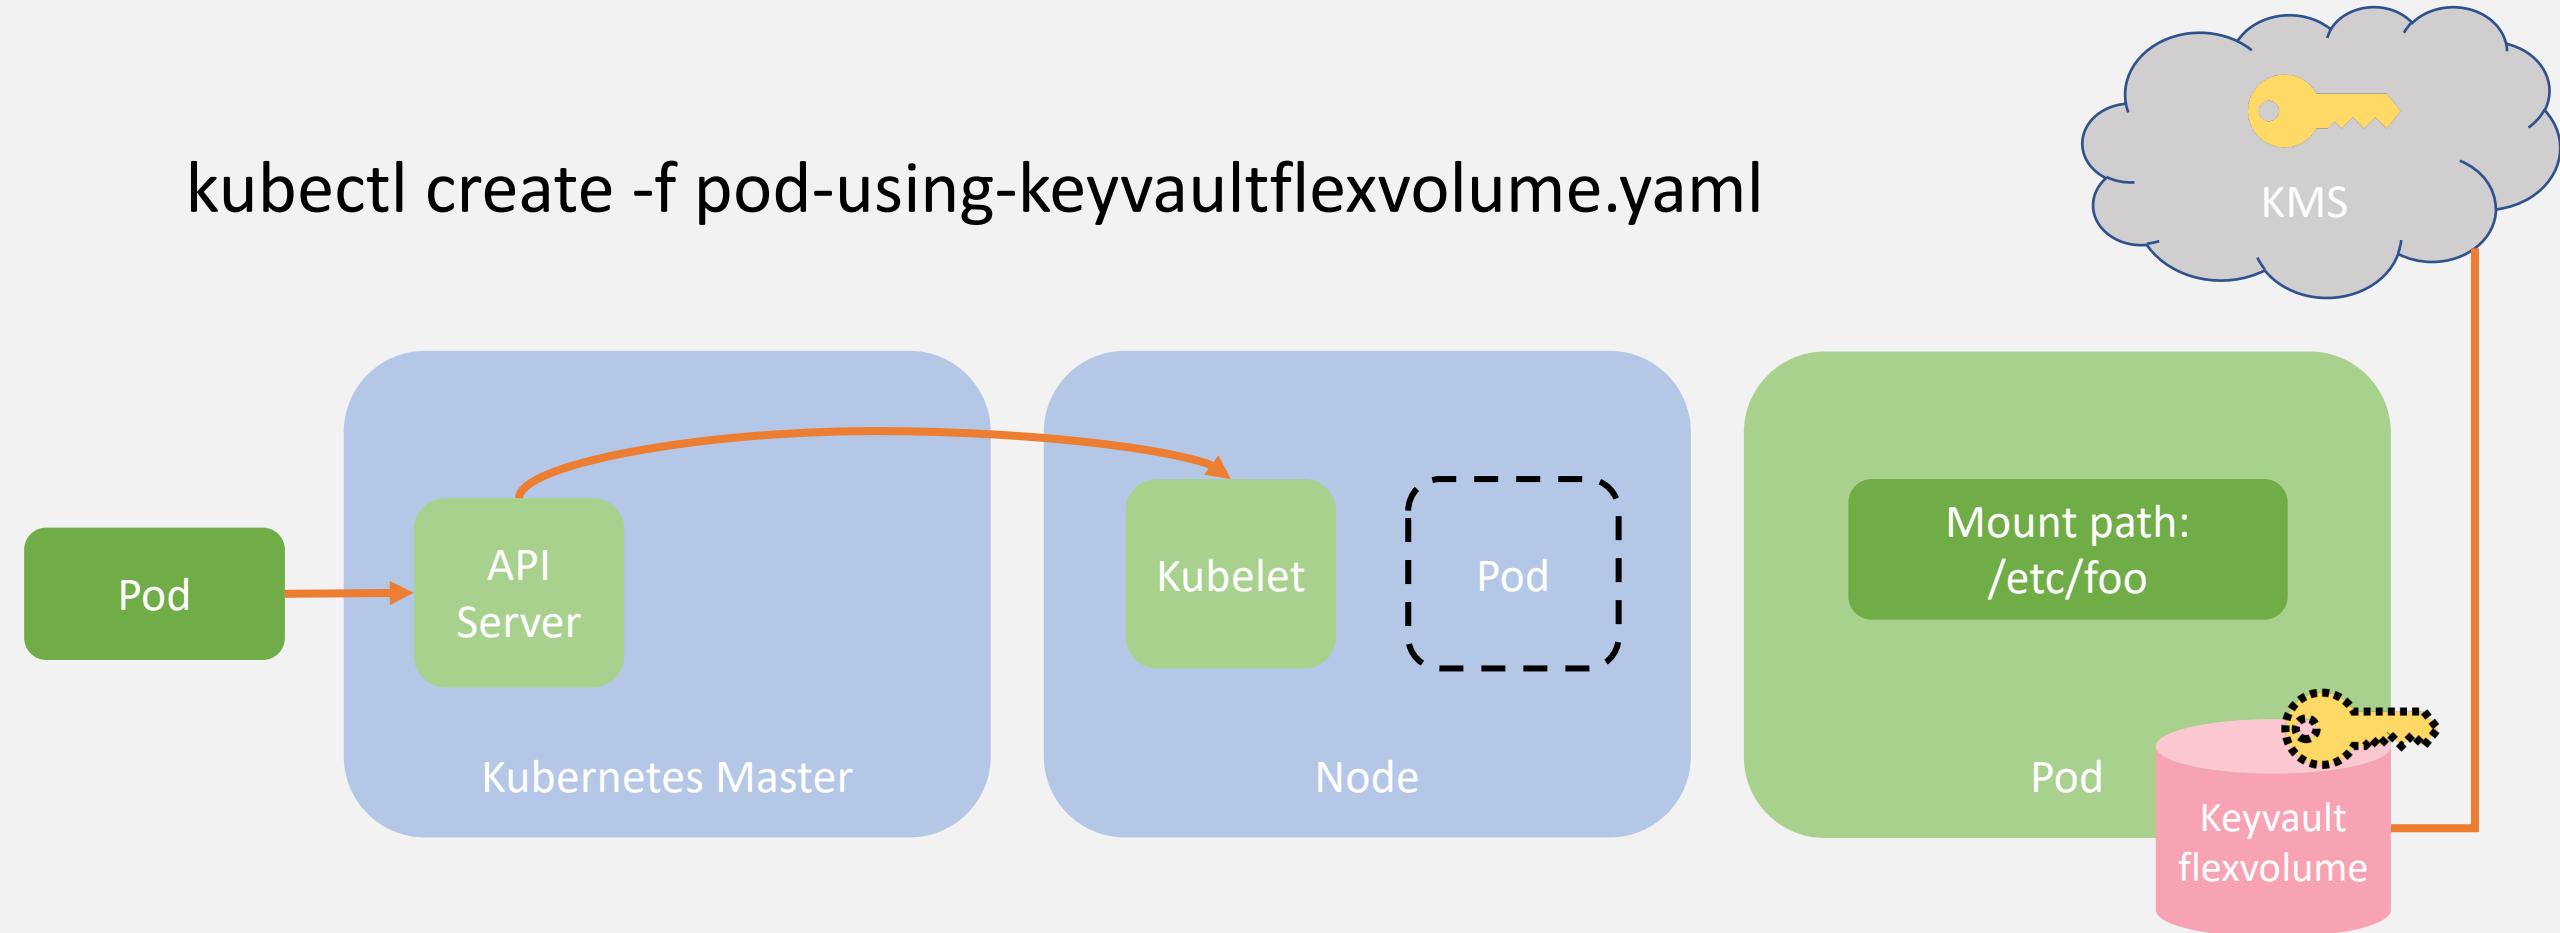


With the KMS provider plugin, we can encrypt Kubernetes data stored in etcd at rest with a KMS managed key.

What if instead of storing my secrets in etcd, I want to store and manage access to application secrets outside of Kubernetes?

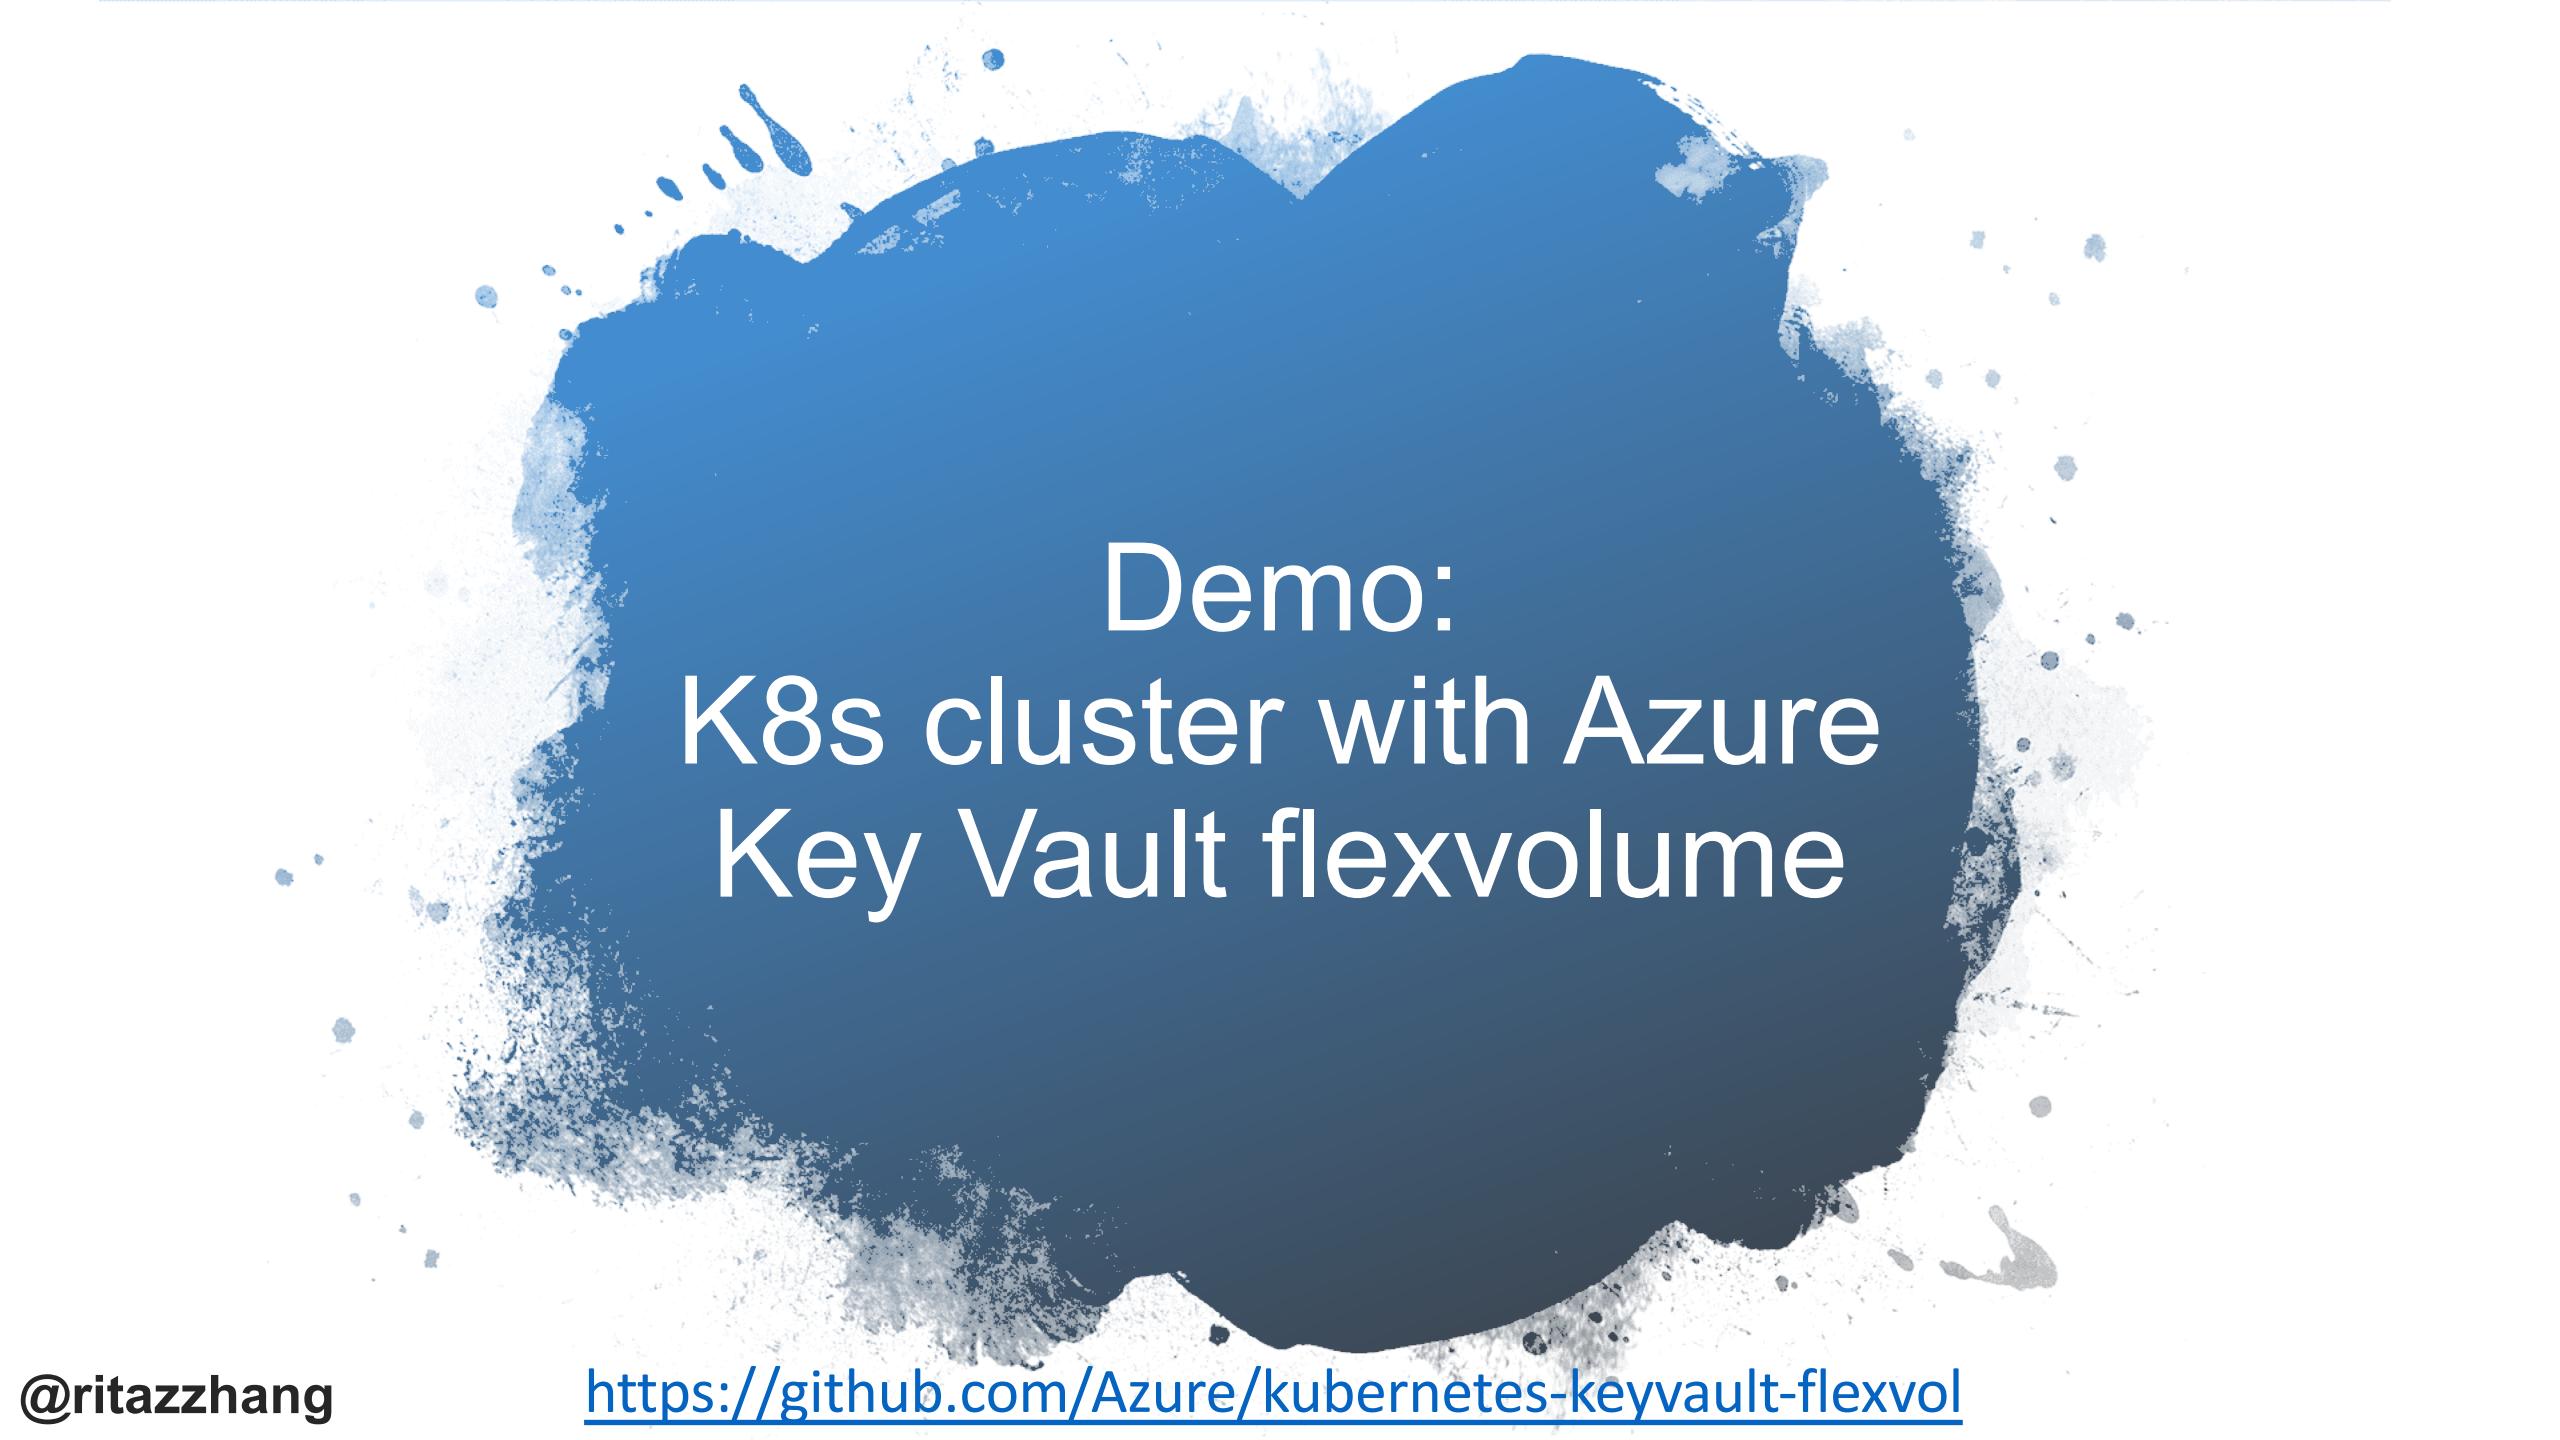
Pod using Key Vault Flexvolume

```
kubectl create -f pod-using-keyvaultflexvolume.yaml
```



Kubernetes Key Vault FlexVolume

- Flexvolume enables users to mount vendor volumes into kubernetes. It expects vendor drivers to be installed in the volume plugin path on every kubelet node.
- Key Vault flexvol driver makes a request to Key Management Service, (e.g. Azure Key Vault) and mounts secret and secret value as a volume to pods
- Separation of concern/role based management access to secrets



Demo: K8s cluster with Azure Key Vault flexvolume

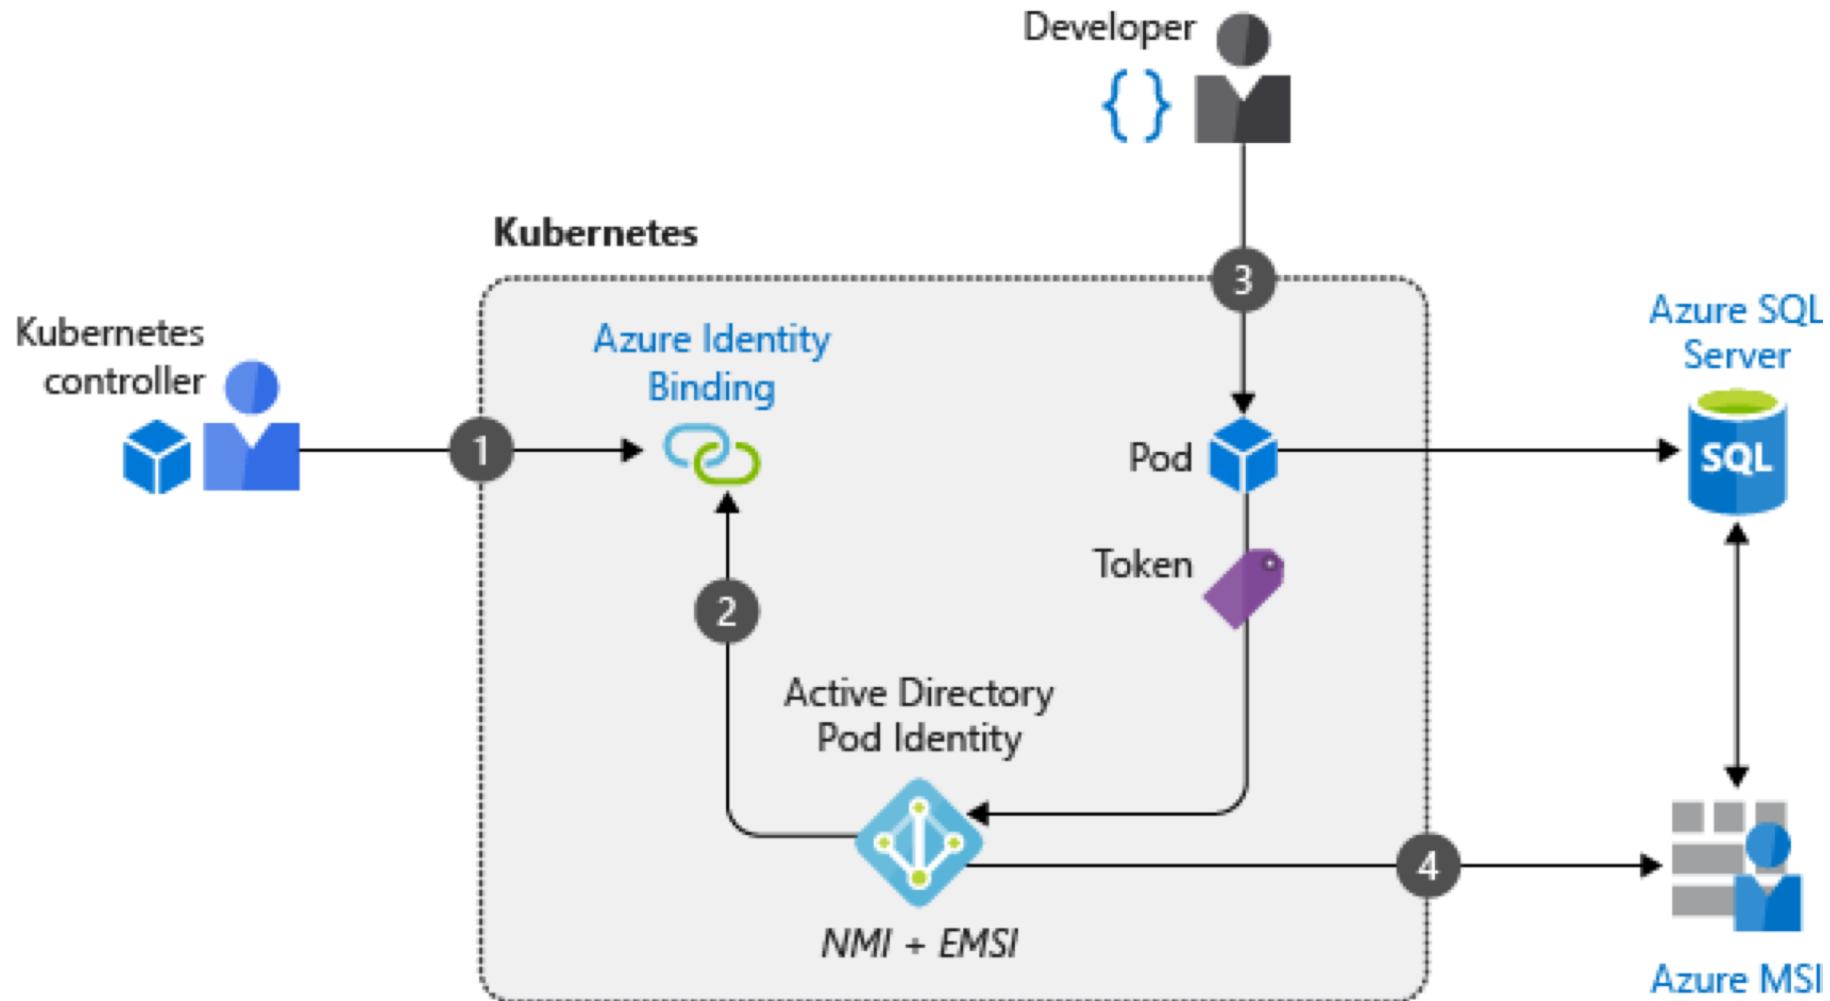
With the Kubernetes Key Vault FlexVolume driver, we can store and retrieve secrets from an Azure Key Vault instance and mount the values as a volume to containers.

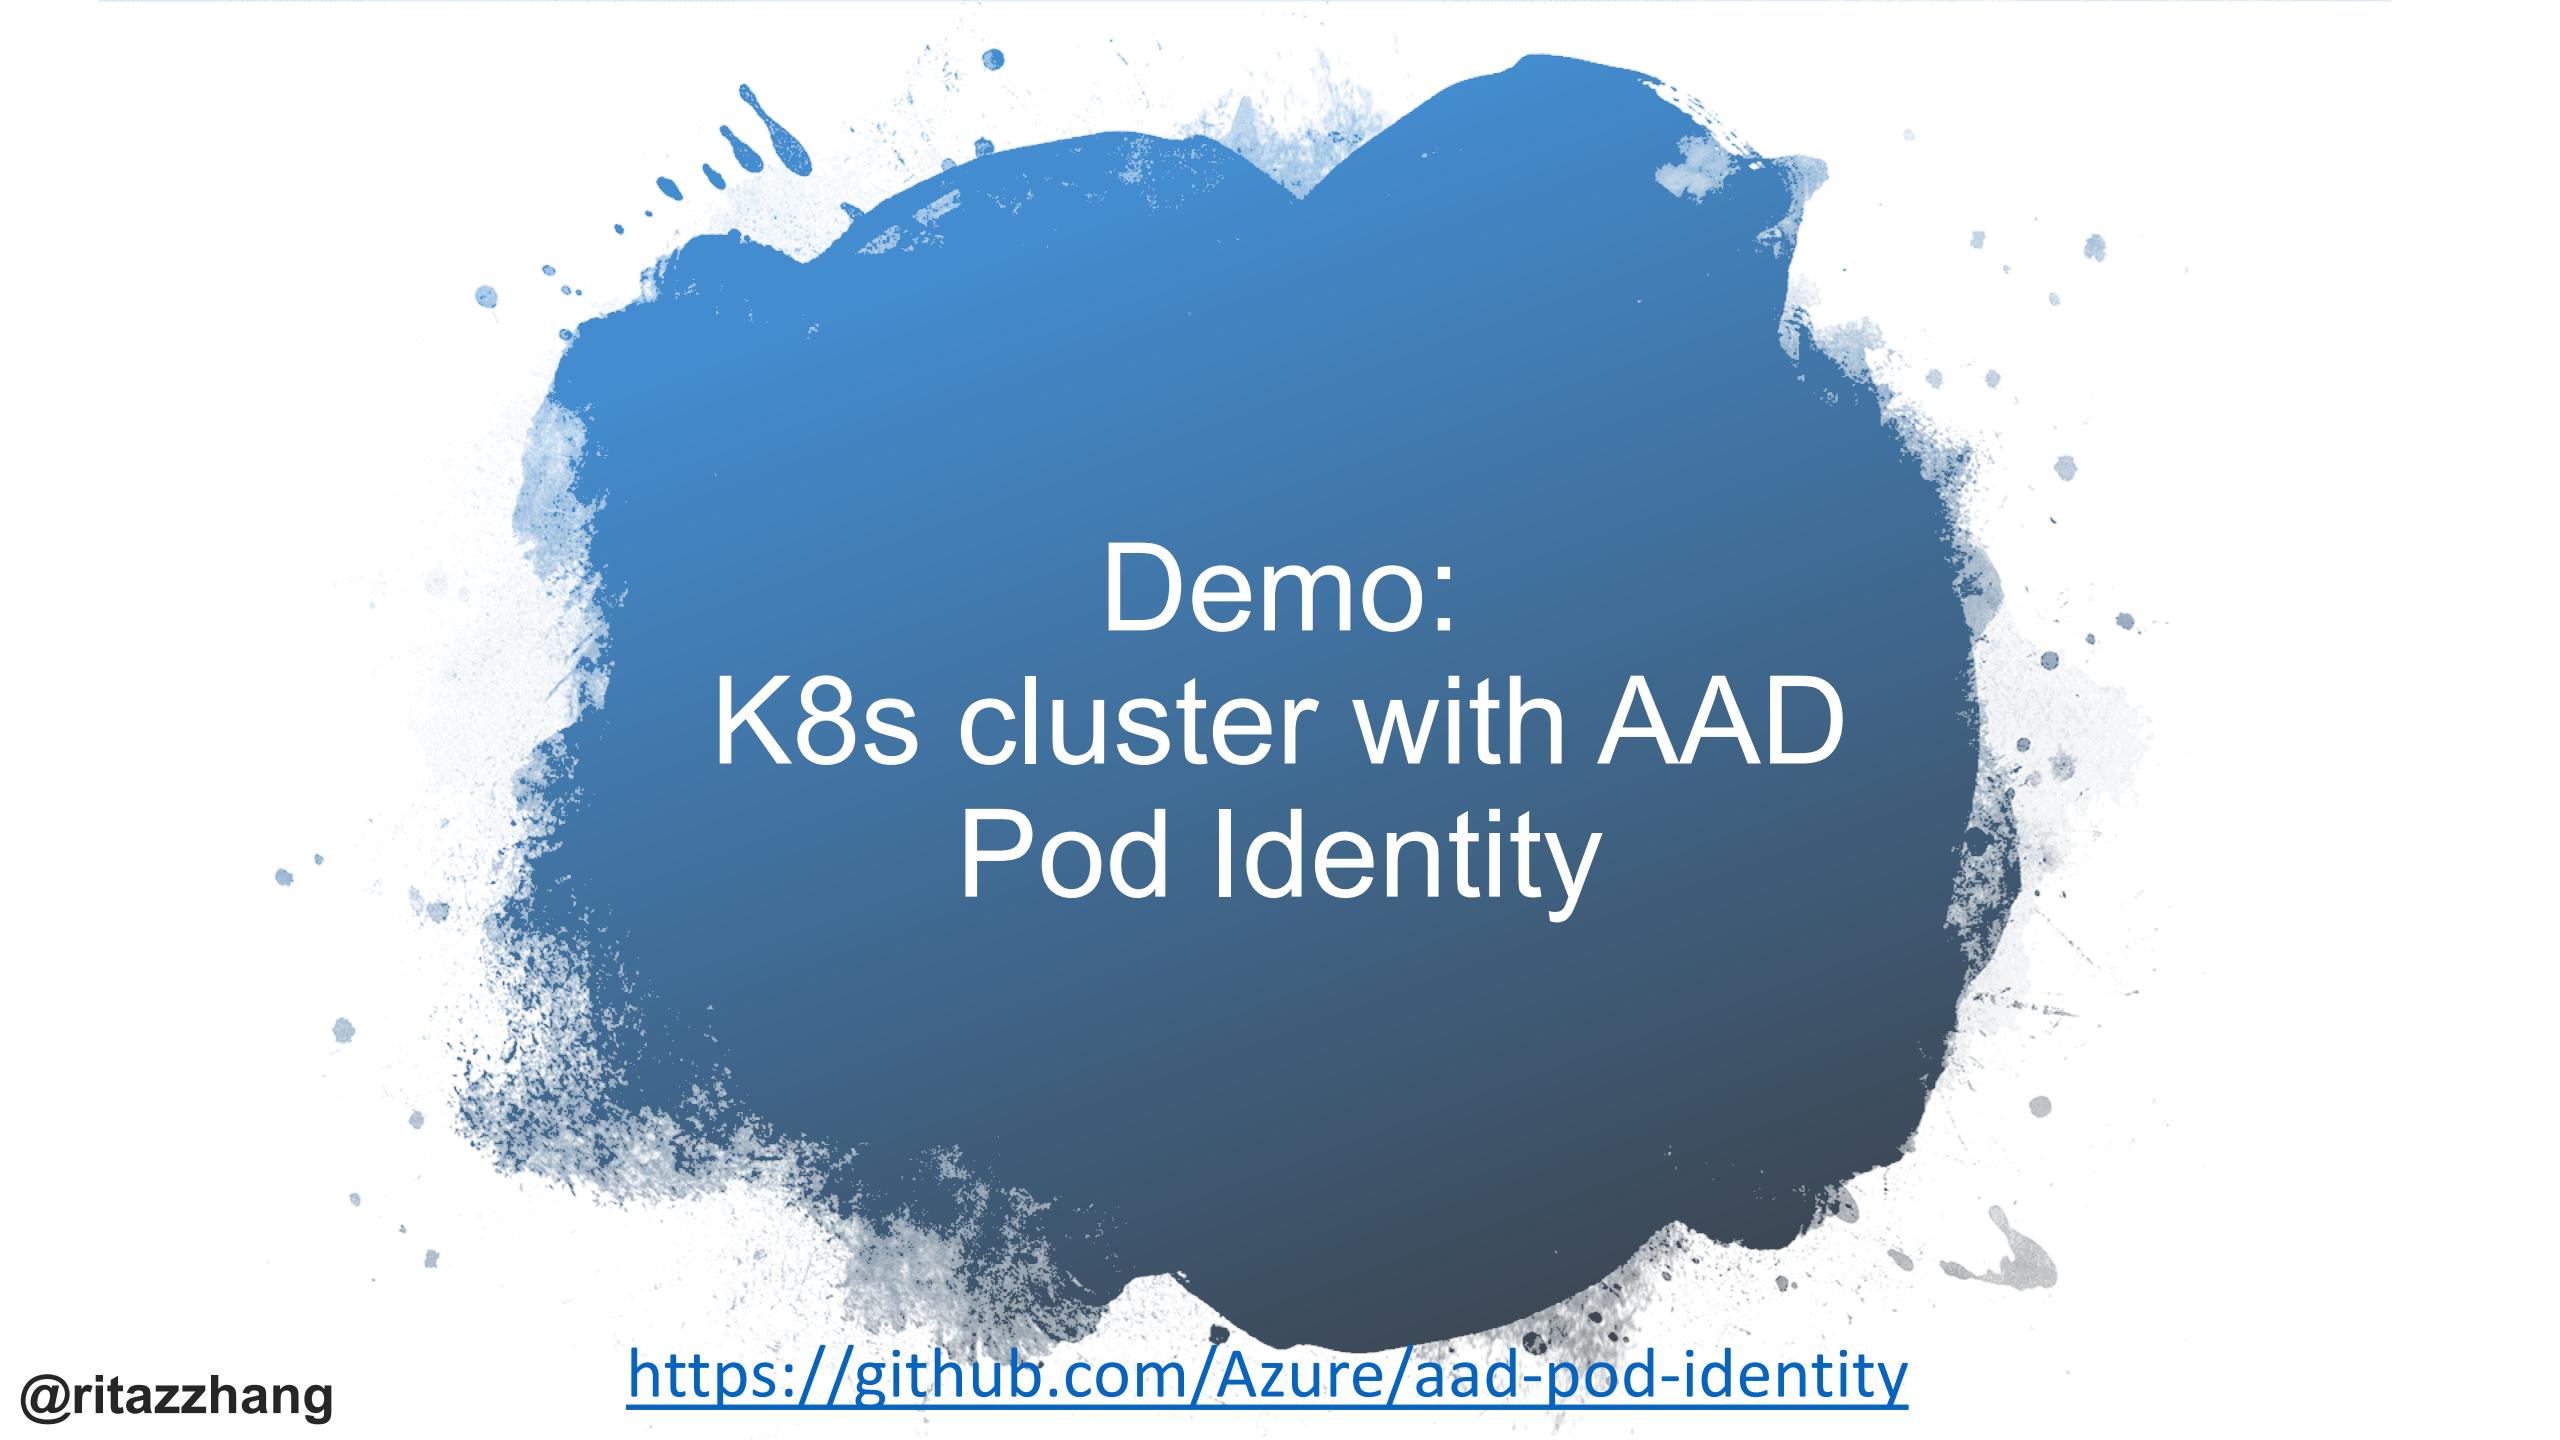
What if I want to restrict access to my cloud resources to specific pods?

AAD Pod Identity

- Restrict/enable pods to access individual resources that depend on Azure AD for access with its own identity (e.g. Azure SQL server or your own custom API that uses AAD)
- Kubernetes Custom Resource Definition objects that map pods to Azure AD identities
- When pods request access to a resource that uses Azure AD for access, a matching Azure identity is assigned

AAD Pod Identity

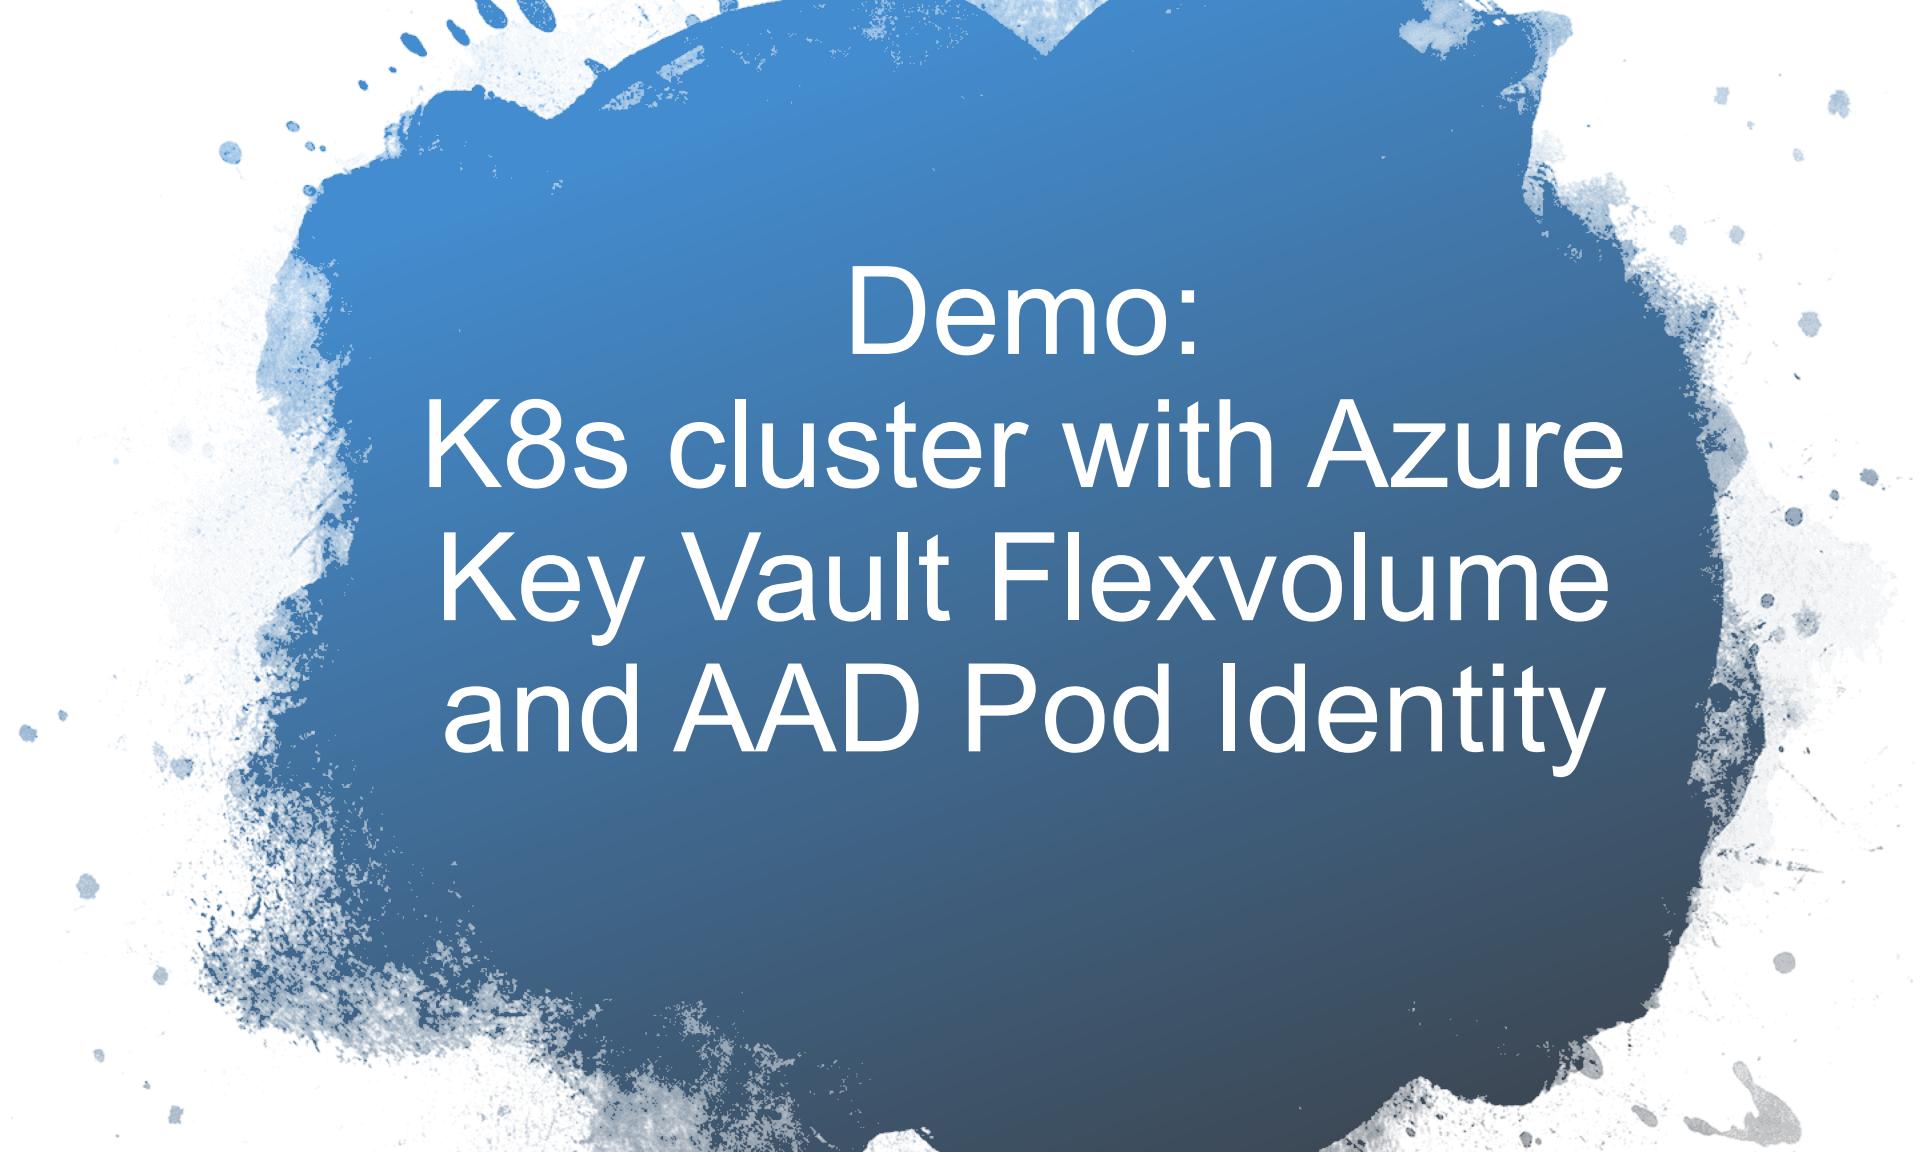




Demo: K8s cluster with AAD Pod Identity

<https://github.com/Azure/aad-pod-identity>

With AAD Pod Identity, we can restrict and enable specific pods access to resources that need Azure AD for access based on its identity.



Demo: K8s cluster with Azure Key Vault Flexvolume and AAD Pod Identity

Recap

Azure Key Vault KMS plugin

- Use a key in Key Vault for etcd encryption
- Secrets/keys/certs are stored in etcd, managed as part of Kubernetes
- Restrict access using K8s concepts: RBAC, Service Accounts, namespaces
- Bring your own keys, import or generate HSM-protected (Hardware Security Modules) keys
- Available on AKS-engine

Flexvol + AAD Pod Identity

- Mounts secrets/keys/certs to the pod using a flexvolume
- Secrets/keys/certs are stored in Azure Key Vault
- Restrict access to secrets/keys/certs with specific pod identities
- Manage secrets in Key Vault via Azure API's/CLI/Portal
- Separation of concerns/role-based management access to secrets
- Import or generate HSM-protected (Hardware Security Modules) keys
- Industry compliance
- More granular RBAC at the pod level

Resources

- Blog post: <https://ritazh.com/using-azure-key-vault-for-kubernetes-data-encryption-d5eac8daee71>
- Kubernetes doc: <https://kubernetes.io/docs/tasks/administer-cluster/kms-provider/>
- aks-engine doc: <https://github.com/Azure/aks-engine/blob/master/docs/kubernetes/features.md#azure-key-vault-data-encryption>
- KMS plugin service PR: <https://github.com/kubernetes/kubernetes/pull/55684>
- Kubernetes KMS Plugin for Azure Key Vault: <https://github.com/Azure/kubernetes-kms>
- Kubernetes KMS Plugin for Google CloudKMS: <https://github.com/GoogleCloudPlatform/k8s-cloudkms-plugin/>
- Kubernetes Key Vault FlexVolume: <https://github.com/Azure/kubernetes-keyvault-flexvol>
- AAD Pod Identity: <https://github.com/Azure/aad-pod-identity>



KubeCon

CloudNativeCon

North America 2018
