

Batch Encoding Audio with Kubernetes

Leigh Capili
Infrastructure Engineer, Beatport

John Slivka
Infrastructure Engineer, Beatport



Leigh Capili

- Infrastructure Engineer
- SIG Cluster-Lifecycle
- @capileigh
- github: stealthybox





John Slivka

- Infrastructure Engineer
- <https://github.com/jslivka>
- @jslivka



GENRES ▾

TRACKS

RELEASES

CHARTS

STEMS

SOUNDS

Q SEARCH...



 My Beatport



NEW ON BEATPORT

MACEO PLEX
MUTANT ROMANCE
MPLX



MPLX

01

EXCLUSIVE

Mutant Romance ^a

Maceo Plex | MPLX



\$4.98



NEW RELEASES



Breathe Original Mix

CamelPhat, Jem Cooke, Cristoph


<https://www.beatport.com/release/mutant-romance/2448894>

TUVABAND VS. RITCHIE TREES


DEAR

125 BPM / A min / 7:13

DJ CHARTS



Maceo Plex



Jack Back



Amelie Lens



beatport top10

1

Hanging Tree Original Mix

Michael Bibi

Repopulate Mars

2

Breathe Original Mix

CamelPhat, Jem Cooke, Cristoph

Pryda Presents

3

Along Came Polly Original Mix

Rebuke

Hot Creations

4

Praise You Purple Disco Machine Extended Remix

Purple Disco Machine, Fatboy Slim

Defected

5

Do It Like This Extended Mix

Biscits

SOLOTOKO

6

Push Pull Club Mix

\$1.99



NEW ON BEATPORT

MACEO PLEX
MUTANT ROMANCE
MPLX



MPLX

01

Mutant Romance ^a

Maceo Plex | MPLX

  \$4.98 ▾

DJ CHARTS



EXCLUSIVE



Amelie Lens

beatport top 10

Hanging Tree Original Mix

[Michael Bibi](#)

Repopulate Mars

Breathe Original Mix

[CamelPhat](#), [Jem Cooke](#), [Cristoph](#)

Pryda Presents

Along Came Polly Original Mix

[Rebuke](#)

Hot Creations

Praise You Purple Disco Machine Extended Remix

[Purple Disco Machine](#), [Fatboy Slim](#)

Defected

5 Do It Like This Extended Mix

[Biscits](#)

SOLOTOKO

6 Push Pull Club Mix

NEW RELEASES



Breathe Original Mix

[CamelPhat](#), [Jem Cooke](#), [Cristoph](#)

<https://www.beatport.com/release/mutant-romance/244RR94>

125 BPM / A min / 7:13

\$1.99 ▾



5

NEW ON BEATPORT

MACEO PLEX
MUTANT ROMANCE
MPLX



MPLX

01

Mutant Roman

Maceo Plex | MPLX



\$4.98



NEW RELEASES



Breathe Original Mix

CamelPhat, Jem Cooke, Cristoph

<https://www.beatport.com/release/mutant-romance/2448894>

EXCLUSIVE



Amelie Lens

DJ CHARTS



beatport top 10

Hanging Tree Original Mix

Michael Bibi

Repopulate Mars

Breathe Original Mix

CamelPhat, Jem Cooke, Cristoph

Pryda Presents

Along Came Polly Original Mix

Rebuke

Hot Creations

Praise You Purple Disco Machine Extended Remix

Purple Disco Machine, Fatboy Slim

Defected

5

Do It Like This Extended Mix

Biscits

SOLO TOKO

6

Push Pull Club Mix

125 BPM / A min / 7:13

\$1.99 ▾



NEW ON BEATPORT

DJ CHARTS

beatport top 10

MACEO PLEX
MUTANT ROMANCE
MPLX



MPLX

01

Mutant Roman

Maceo Plex | MPLX



\$4.98

NEW RELEASES

Breathe Original Mix

CamelPhat, Jem Cooke, Cristoph

<https://www.beatport.com/release/mutant-romance/2448894>



EXCLUSIVE



Hanging Tree Original Mix

Michael Bibi

Repopulate Mars

Breathe Original Mix

CamelPhat, Jem Cooke, Cristoph

Pryda Presents

Along Came Polly Original Mix

Rebuke

Hot Creations

Praise You Purple Disco Machine Extended Remix

Purple Disco Machine, Fatboy Slim

Defected

5 Do It Like This Extended Mix

Biscits

SOLOTOKO

6 Push Pull Club Mix

125 BPM / A min / 7:13

\$1.99 ▾



7



GENRES ▾

TRACKS

RELEASES

CHARTS

STEMS

SOUNDS

Q SEARCH...



My Beatport



NEW ON BEATPORT

DJ CHARTS



beatport top10

MACEO PLEX
MUTANT ROMANCE
MPLX



MPLX

01

Mutant Roman

Maceo Plex | MPLX



\$4.98

NEW RELEASES

Breathe Original Mix

CamelPhat, Jem Cooke, Cristoph

<https://www.beatport.com/release/mutant-romance/2448894>



Hanging Tree Original Mix

Michael Bibi

Repopulate Mars

Breathe Original Mix

CamelPhat, Jem Cooke, Cristoph

Pryda Presents

Along Came The Original Mix

Re

Extended Remix


Push Pull Club Mix

125 BPM / A min / 7:13




\$1.99 ▾



8


GENRES ▾TRACKSRELEASESCHARTSSTEMSSOUNDS


Q SEARCH...



 My Beatport 

NEW ON BEATPORT


MACEO PLEX
MUTANT ROMANCE
MPLX



5 060587 486133 >
MPLX
01





Mutant Romance
Maceo Plex | MPLX
  **\$4.98** ▾

NEW RELEASES

**Breathe** Original Mix
CamelPhat, Jem Cooke, Cristoph
<https://www.beatport.com/release/mutant-romance/2448894>



125 BPM / A min / 7:13

  **beatport** top10

1 **Hanging Tree** Original Mix
Michael Bibi
Repopulate Mars









2 **Breathe** Original Mix
CamelPhat, Jem Cooke, Cristoph
Pryda Presents

3 **Along Came Polly** Original Mix
Rebuke
Hot Creations

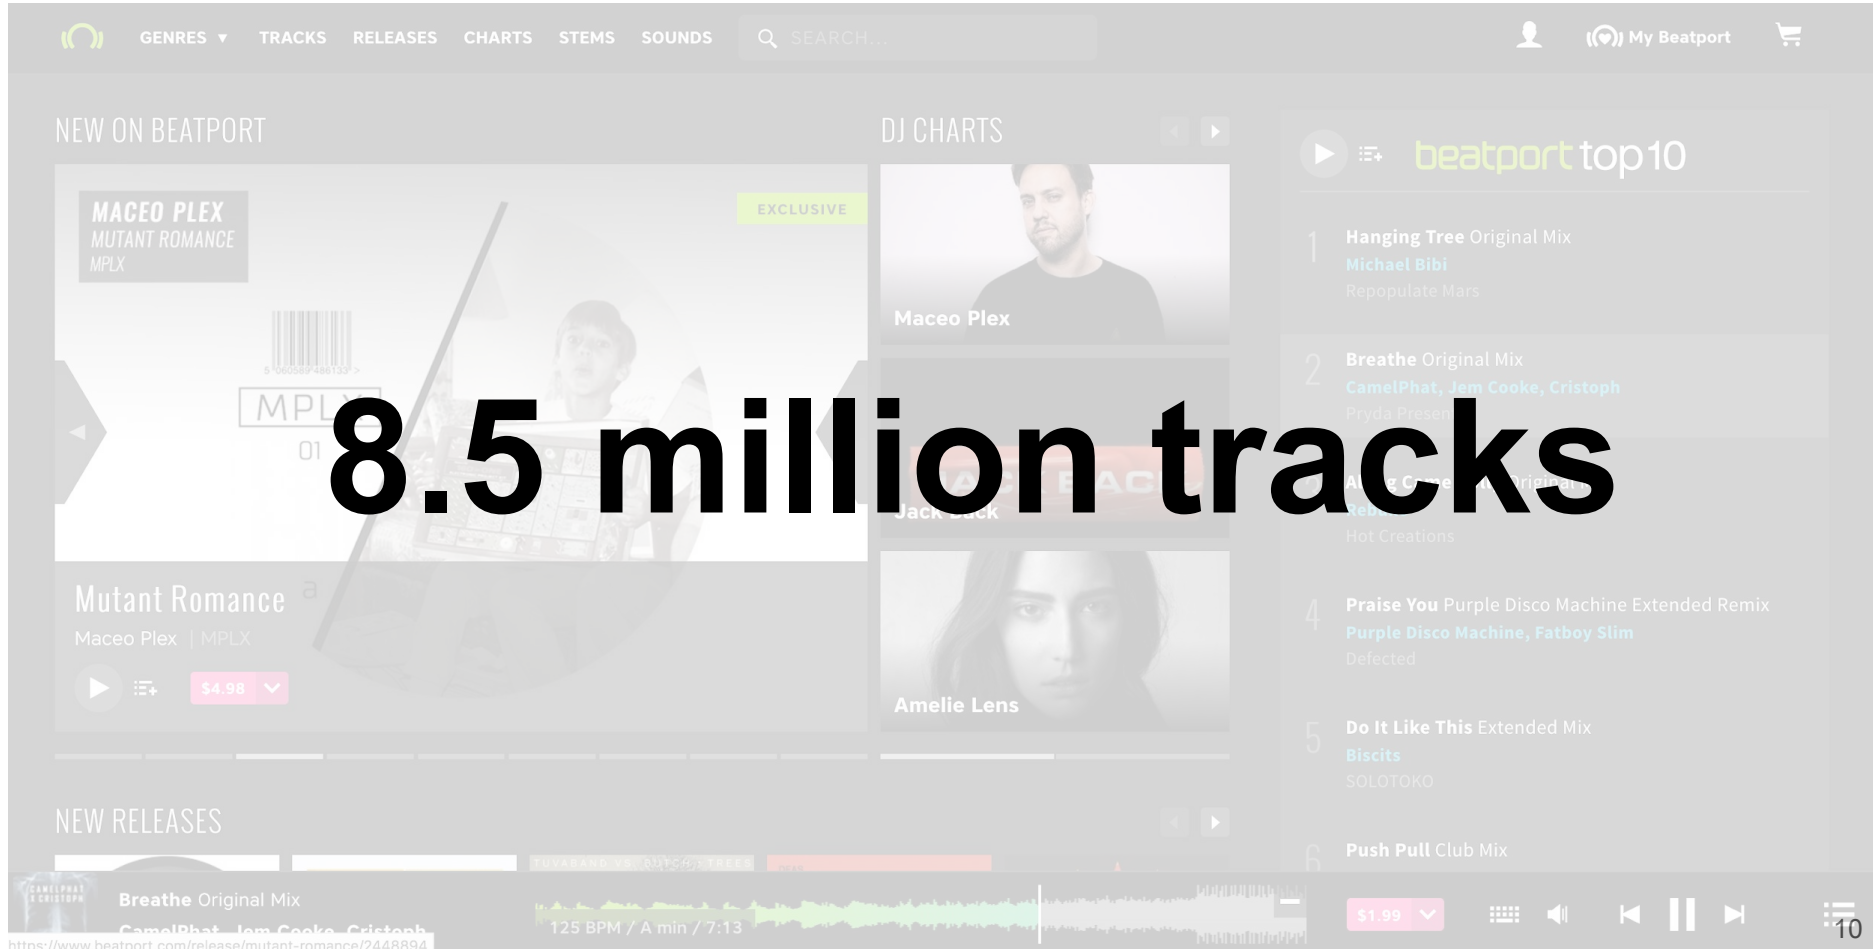
4 **Praise You** Purple Disco Machine Extended Remix
Purple Disco Machine, Fatboy Slim
Defected

5 **Do It Like This** Extended Mix
Biscits
SOLOTOKO

6 **Push Pull** Club Mix

  **\$1.99** ▾      

9



The image is a screenshot of the Beatport website. At the top, there is a navigation bar with the Beatport logo, links for GENRES, TRACKS, RELEASES, CHARTS, STEMS, and SOUNDS, a search bar, and user account options. The main content area is divided into several sections: 'NEW ON BEATPORT' featuring a large promotional banner for 'MACEO PLEX MUTANT ROMANCE MPLX' with a barcode and a play button; 'DJ CHARTS' showing a grid of artist portraits including Maceo Plex, Jack Back, and Amelie Lens; and a 'beatport top 10' list on the right. The top 10 list includes tracks like 'Hanging Tree Original Mix' by Michael Bibi, 'Breathe Original Mix' by CamelPhat, Jem Cooke, and Cristoph, 'Praise You Purple Disco Machine Extended Remix' by Purple Disco Machine and Fatboy Slim, 'Do It Like This Extended Mix' by Biscits, and 'Push Pull Club Mix'. At the bottom, there is a 'NEW RELEASES' section and a music player interface showing the track 'Breathe Original Mix' by CamelPhat, Jem Cooke, and Cristoph, with a waveform and playback controls. A large, bold, black text overlay '8.5 million tracks' is centered across the middle of the page.

8.5 million tracks

Agenda

- Timeline
- Business Goal
- Problem Definition
- Approach
- Prototype
- Performance
- Scheduling
- Operations
- Conclusion
- Q&A

Business Goal

- Batch process our entire back catalog to-date (8-9 million tracks) to render the following derivative assets
 - Recalculate BPM (more accurate algorithm)
 - Recalculate Key (substantially more accurate algorithm)
 - 128K AAC for full-length preview and download
 - Capability to plug-in additional derivative assets into the system easily with minimal/no downtime

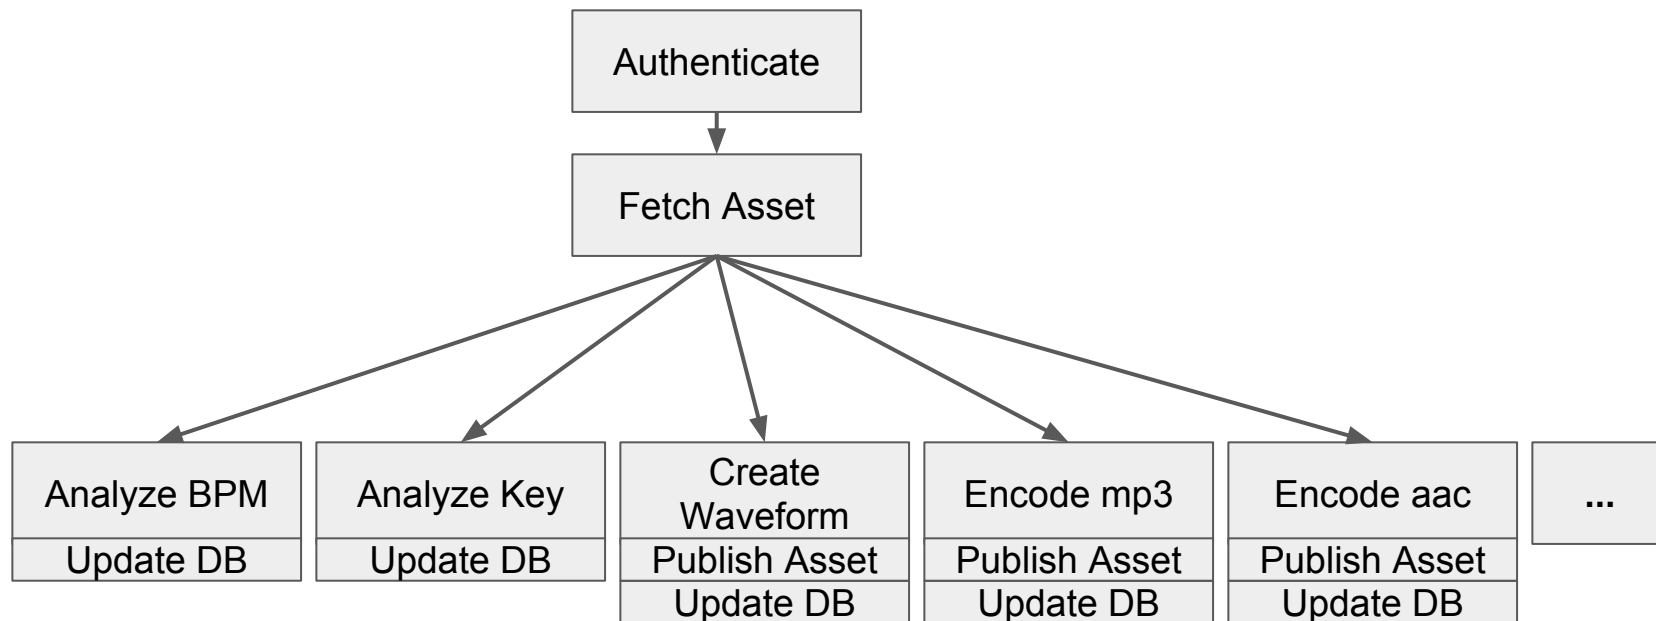
Terms

- **Master:** authoritative source of audio provided by our suppliers to be sold in the store. 16-bit stereo 44.1K PCM data stored as WAV
- **Derivative Asset:** Any information that is *derived* from an audio master (lossless and lossy compressions, waveforms, metadata like BPM and Key)
- **Back Catalog:** Collection of masters we keep that are currently available in our store for purchase
- **Release:** Single, EP or LP. Collection of content including tracks. Also: album artwork and additional metadata
- **Task:** collection of work to produce derivative assets for a single track (ultimately they share a single pod)

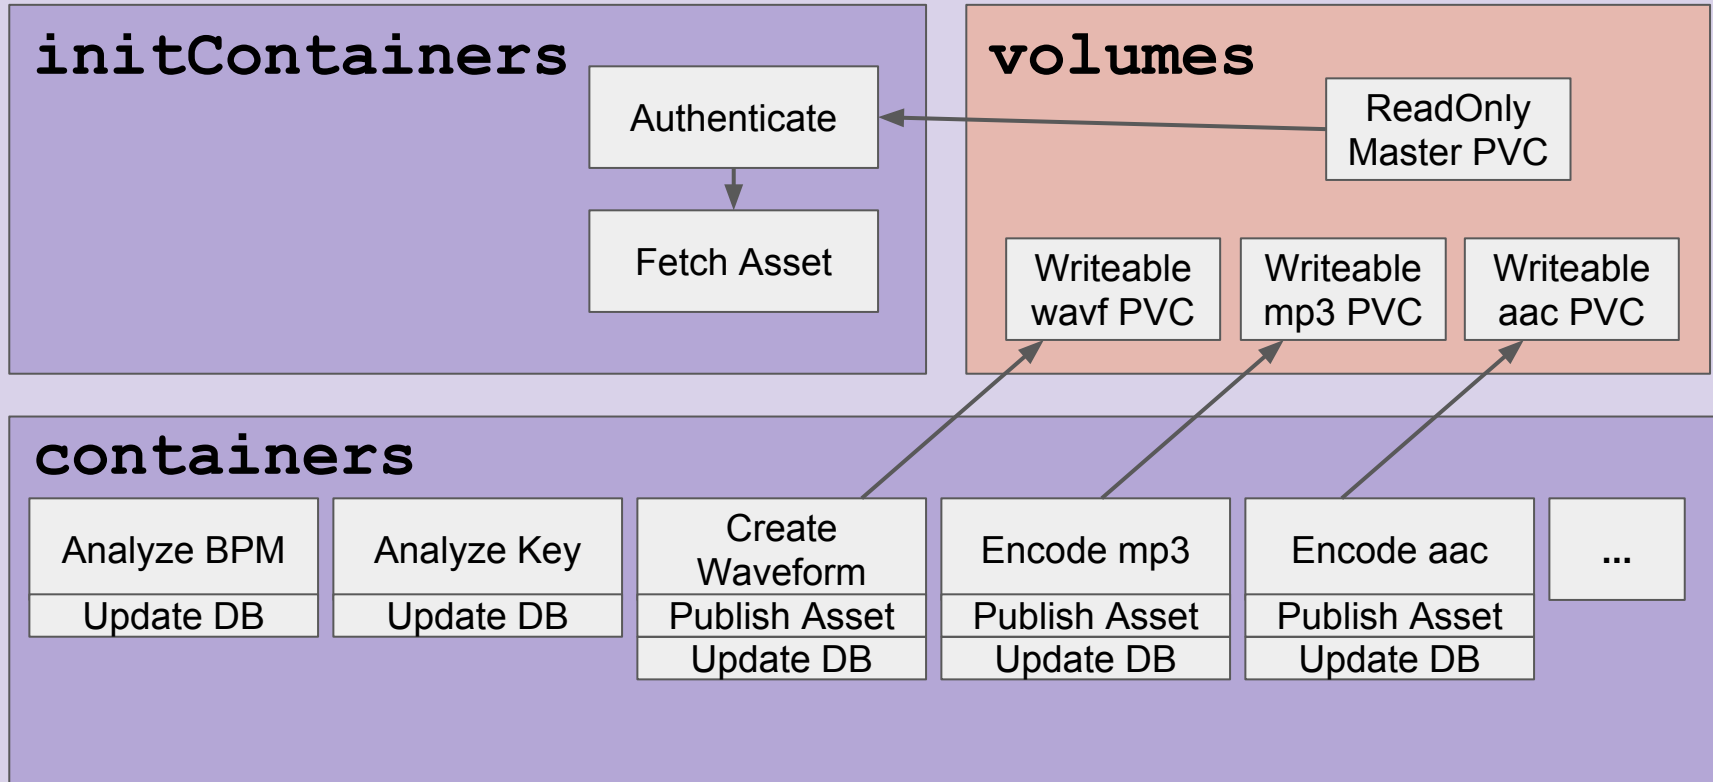
Constraints / Goals

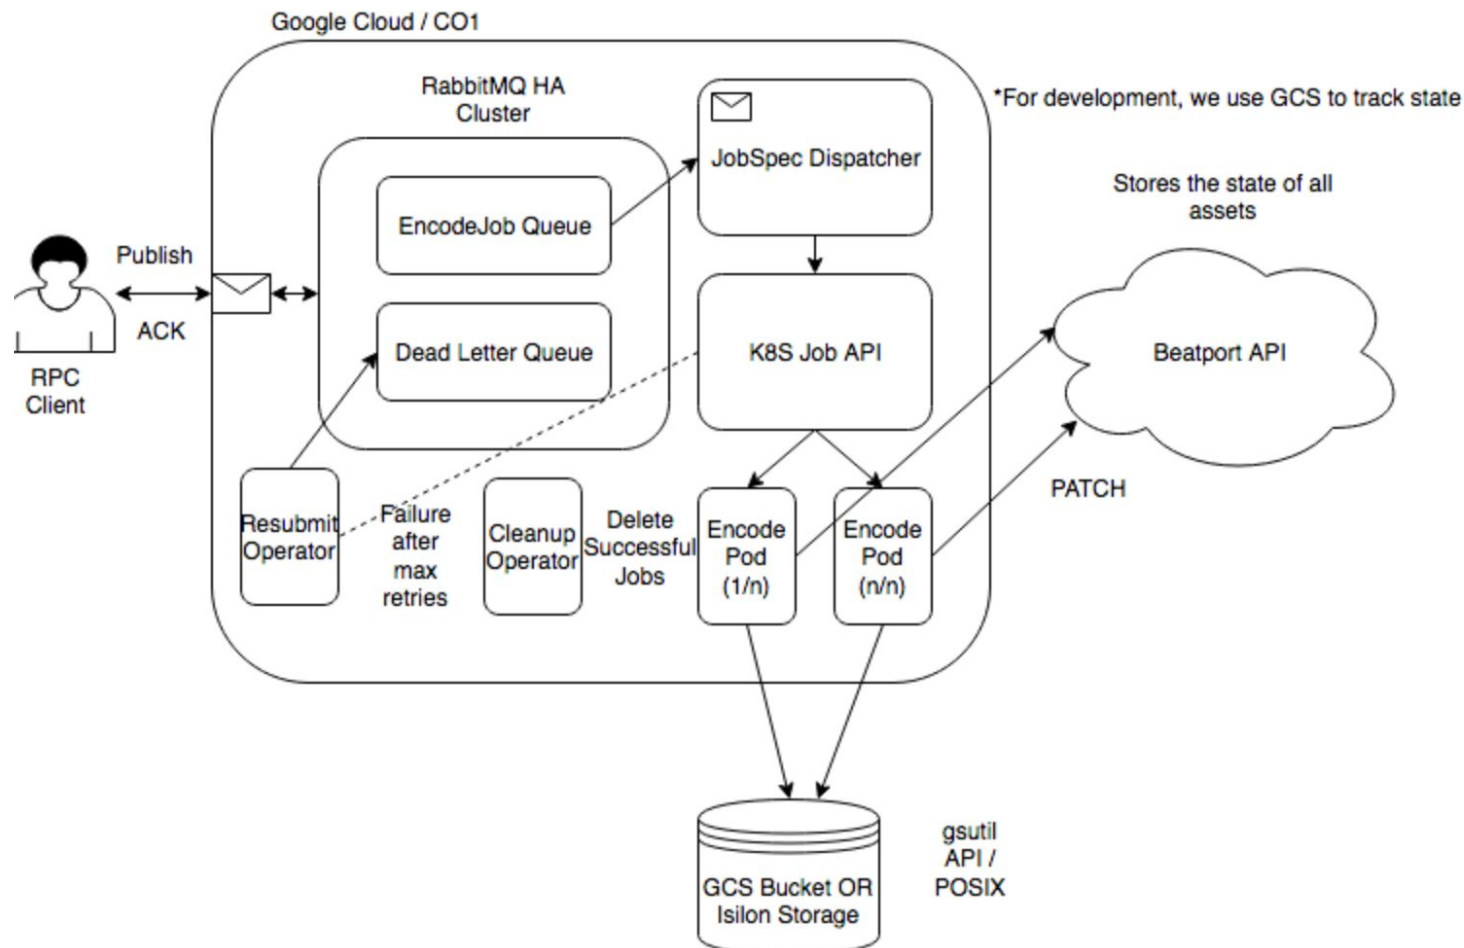
- Newer products running on Google Kubernetes Engine (GKE)
- Old enterprise hardware in our datacenter
- Portability
- Storage Models: NFS, Google Cloud Storage Buckets
- Cross-WAN connection from datacenter to GKE, Minimize round-trips
- Observability more challenging for ephemeral workloads
- Handling concurrency and maximizing completion rate is tricky

Business Logic



Pod Representation





What information
does the system
need?

```
{
  "$schema": "http://json-schema.org/draft-06/schema#",
  "title": "EncodeJob",
  "description": "Contract to perform rendering of a derivative asset",
  "type": "object",
  "properties": {
    "request_id": {
      "description": "UUID for a client request",
      "type": "string"
    },
    "release_id": {
      "description": "Regular id for a release from a vendor",
      "type": "integer"
    },
    "asset_id": {
      "description": "Track ID",
      "type": "integer"
    },
    "asset_guid": {
      "description": "Track GUID",
      "type": "string"
    },
    "derivative_assets": {
      "type": "array",
      "items": {
        "type": "string"
      },
      "uniqueItems": true
    }
  }
}
```

This will mostly be used for processing release artwork (image resizing)



```
{
  "$schema": "http://json-schema.org/draft-06/schema#",
  "title": "EncodeJob",
  "description": "Contract to perform rendering of a derivative asset",
  "type": "object",
  "properties": {
    "request_id": {
      "description": "UUID for a client request",
      "type": "string"
    },
    "release_id": {
      "description": "Regular id for a release from a vendor",
      "type": "integer"
    },
    "asset_id": {
      "description": "Track ID",
      "type": "integer"
    },
    "asset_guid": {
      "description": "Track GUID",
      "type": "string"
    },
    "derivative_assets": {
      "type": "array",
      "items": {
        "type": "string"
      },
      "uniqueItems": true
    }
  }
}
```

Integer ID we use
to reference asset
in catalog API



```
{
  "$schema": "http://json-schema.org/draft-06/schema#",
  "title": "EncodeJob",
  "description": "Contract to perform rendering of a derivative asset",
  "type": "object",
  "properties": {
    "request_id": {
      "description": "UUID for a client request",
      "type": "string"
    },
    "release_id": {
      "description": "Regular id for a release from a vendor",
      "type": "integer"
    },
    "asset_id": {
      "description": "Track ID",
      "type": "integer"
    },
    "asset_guid": {
      "description": "Track GUID",
      "type": "string"
    },
    "derivative_assets": {
      "type": "array",
      "items": {
        "type": "string"
      },
      "uniqueItems": true
    }
  }
}
```


Used for storage
of an asset.
Introduced
recently so we
must include this
as well



```
{
  "$schema": "http://json-schema.org/draft-06/schema#",
  "title": "EncodeJob",
  "description": "Contract to perform rendering of a derivative asset",
  "type": "object",
  "properties": {
    "request_id": {
      "description": "UUID for a client request",
      "type": "string"
    },
    "release_id": {
      "description": "Regular id for a release from a vendor",
      "type": "integer"
    },
    "asset_id": {
      "description": "Track ID",
      "type": "integer"
    },
    "asset_guid": {
      "description": "Track GUID",
      "type": "string"
    },
    "derivative_assets": {
      "type": "array",
      "items": {
        "type": "string"
      },
      "uniqueItems": true
    }
  }
}
```

Give me a list of
things to produce

```
{
  "$schema": "http://json-schema.org/draft-06/schema#",
  "title": "EncodeJob",
  "description": "Contract to perform rendering of a derivative asset",
  "type": "object",
  "properties": {
    "request_id": {
      "description": "UUID for a client request",
      "type": "string"
    },
    "release_id": {
      "description": "Regular id for a release from a vendor",
      "type": "integer"
    },
    "asset_id": {
      "description": "Track ID",
      "type": "integer"
    },
    "asset_guid": {
      "description": "Track GUID",
      "type": "string"
    },
    "derivative_assets": {
      "type": "array",
      "items": {
        "type": "string"
      },
      "uniqueItems": true
    }
  }
}
```



Prototype

- Need to regulate flow of job requests to k8s API: use a queue
- Decided against single-purpose daemons so we only fetch an asset once, don't need complex scheduling logic
- Declare a message spec
- Send to queue
- A “dispatcher” will interpret the sent message and mint a jobspec (or podspec) to send to the kubernetes API

Prototype

Good:

- Quick turn-around time to validating idea. (1 week!)
- Pre-populated content samples into a GCS bucket
- GKE allowed us to prototype the software's behavior at scale easily
- Observed the system was behaving with **hjacobs/kube-ops-view**

Frustrating:

- Python kubernetes client was easy to iterate with, but it was difficult to work with type mismatch errors.

Demo :)

Hardware Assets

- 1 Rack
- 12 Dell R520 (2012)
- 1 Dell R410 (2009)
- 1 Dell R620 (2012)
- 4 Dell PE2950 (2006-2008)

Kubernetes v1.11.0


Docker 18.03.1-ce


CentOS 7

Community maintained Linux Kernel (4.16)

Performance / KPIs

- Failure Rate
- Success Rate
- Completions per minute
- Concurrency
- Garbage collection of Tasks
- Requeuing and recording of failed Tasks





Scheduling

- CPU Pinning
- Temporal Slicing
- Requests
- Limits
- <https://hackernoon.com/job-concurrency-in-kubernetes-lxd-and-cpu-pinning-to-the-rescue-b9fb7b44f99d>

A Problem With Jobs

- Kubernetes API had a saturation limit
- Past a certain threshold, CRUD operations became impossibly slow over time (particularly batch deletes)
- Bottlenecked not by CPU, but now by the API

Jobs

- Parallelism
- Retries
- Completions
- “ttlSecondsAfterFinished” automatic GC as of 1.12 alpha

Jobs

- ~~Parallelism~~
- Retries
- ~~Completions~~
- “ttlSecondsAfterFinished” automatic GC as of 1.12 alpha

Pods

- CronJob for garbage collection (Go K8S client)
- CronJob for dead-lettering failed tasks (Go K8S client)

Pods

- Free retries: already have a dead-lettering mechanism built into the queue, these can be triaged shoveled easily
- Acceptable performance
- Still a threshold for objects that should be in the API at our scale
- Change limits on max pods allowed in API

<Placeholder for Graphs>

Finalizing reporting

Conclusion

Questions