



Running Envoy as an Edge Proxy @ eBay

Qiu Yu (QY)
Bala Madhavan

Agenda

- eBay's Edge
- SW Architecture of Edge proxy
 - L4 Load Balancer
 - Managing Envoy in Kubernetes
 - Data Plane Journey
 - Caching on the Edge with Envoy
- Observability
- Deploy management
- Learnings and Challenges
- Q&A

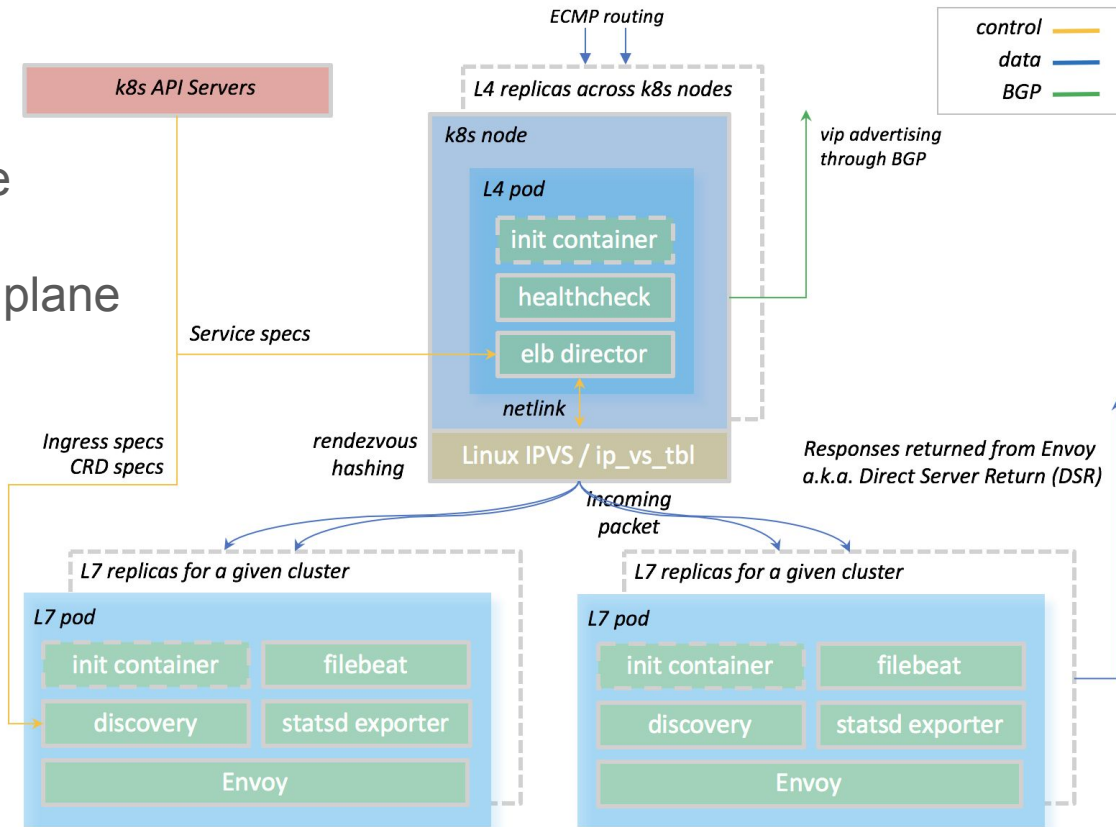
eBay's Edge

- Why Edge
 - Data Centers in US
 - Getting closer to the end user
- Edge PoPs
 - Point of Presences around the Globe
 - Connect eBay network to the Internet
 - Software powered infrastructure
 - k8s cluster in each PoP
- Beyond Speed



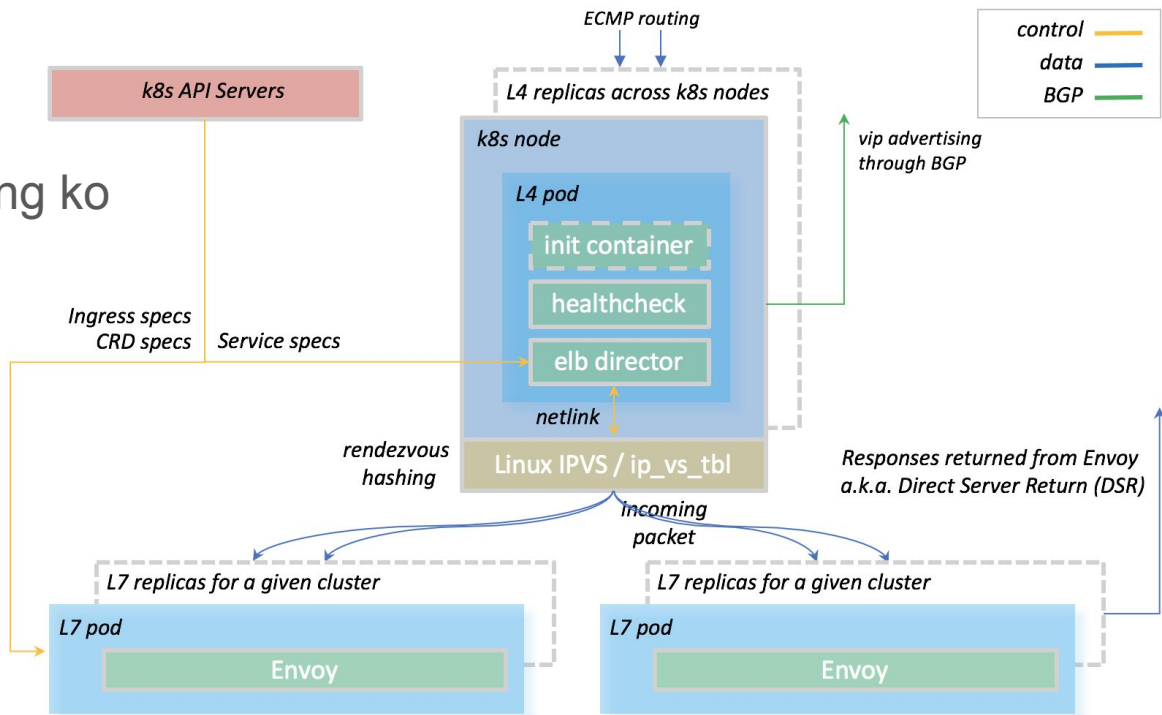
Edge Proxy Architecture

- North-South gateway on Edge
- L4 + L7: Control plane + data plane
- Specs realized by controllers
- K8s native deployment
- Horizontally scalable



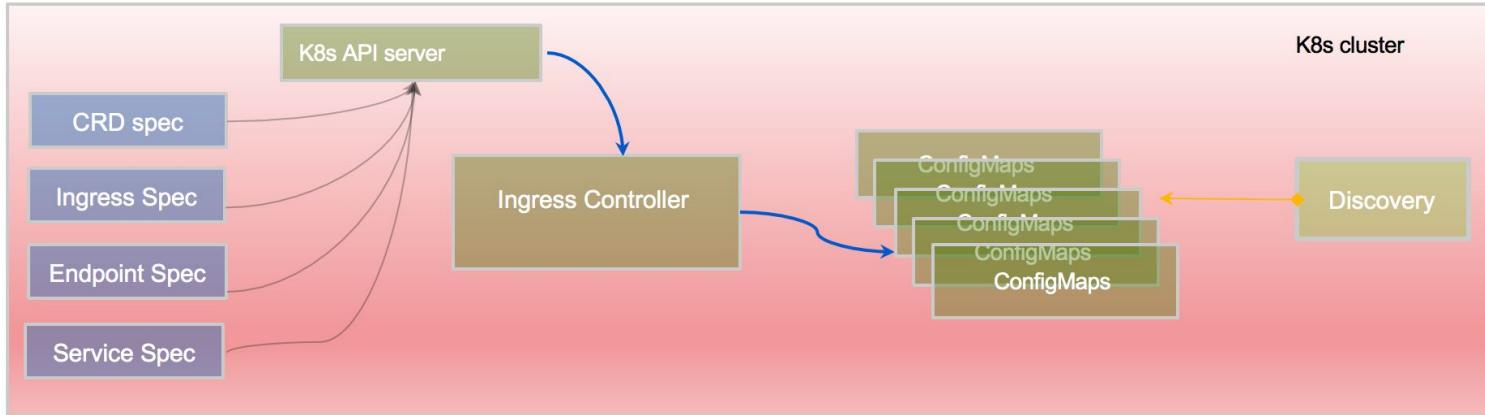
L4 Load Balancer

- IPVS based data plane
- In-house rendezvous hashing ko
- BGP for VIP advertising
- Direct Server Return (DSR)

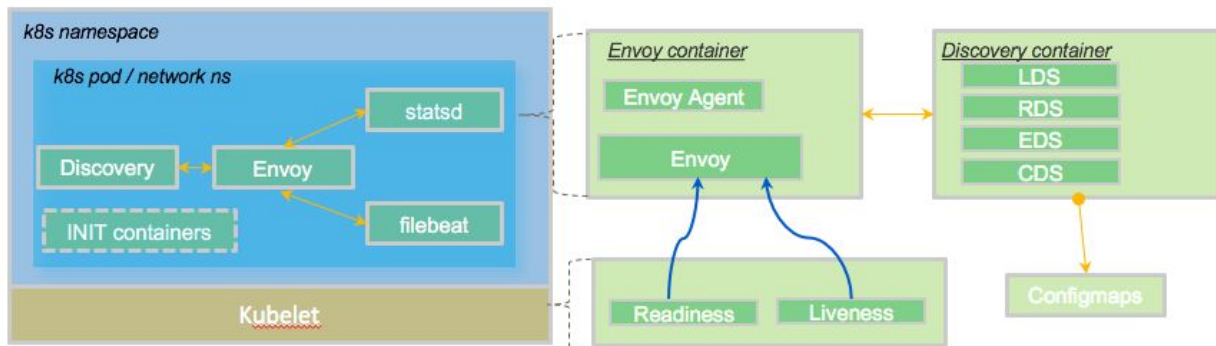


Managing Envoy in Kubernetes

- Workloads: L7 PODs (envoy runs as container) & Deployment controller
 - Rolling updates
 - Scaling
 - Better control (Rollback, Pause, Live Status etc.)
- Listener configured via Ingress spec
 - Annotations for custom features specifications.
- L7 routes managed by Custom Resource Definition (CRDs)



Managing Envoy in Kubernetes - Cont.

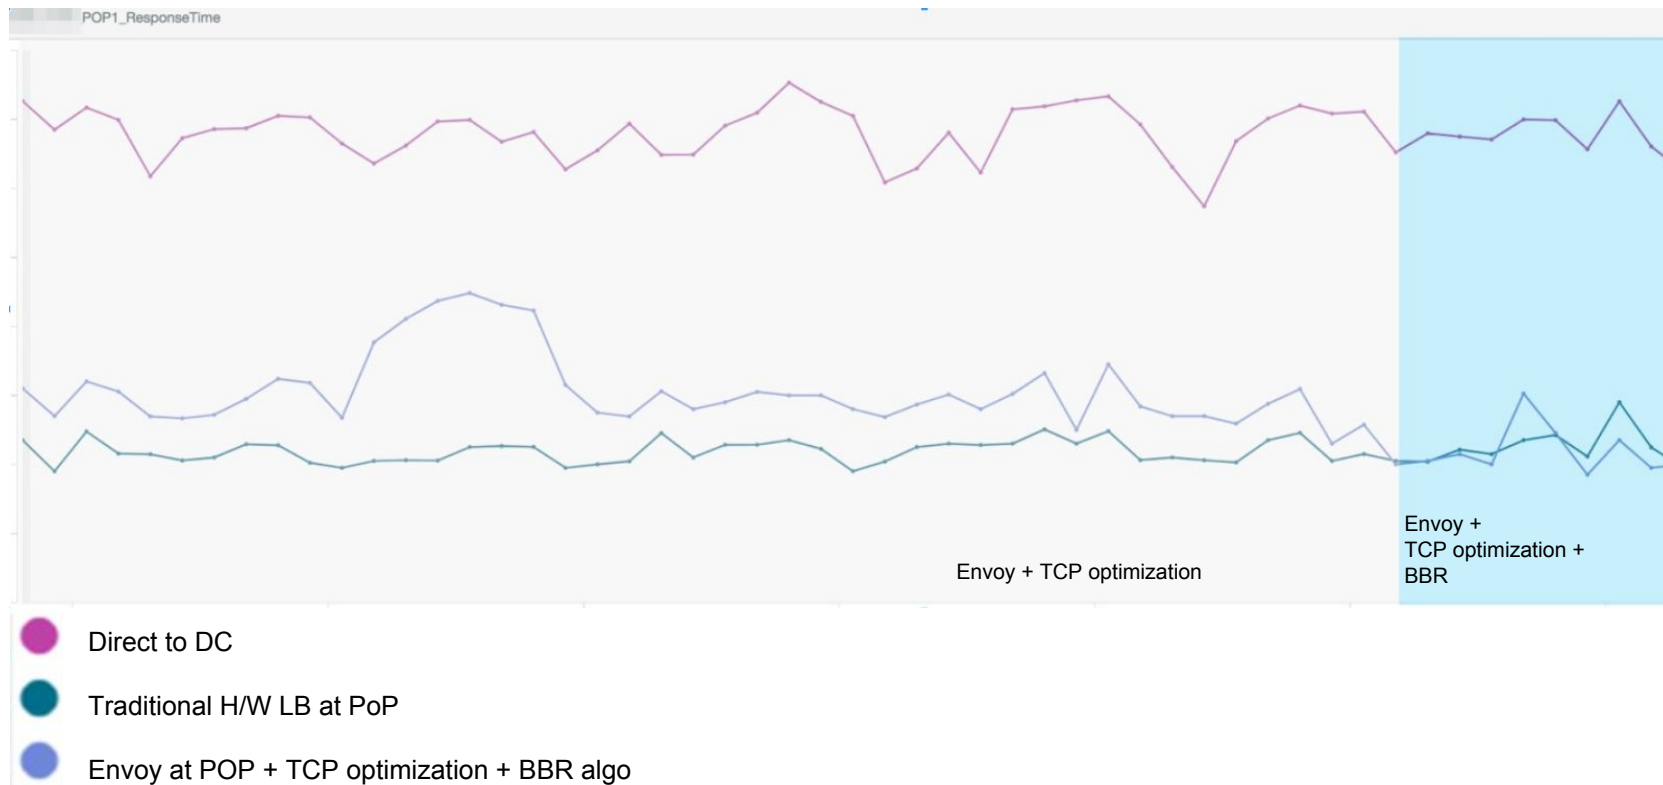


- Custom 'discovery' service handles RDS, LDS, EDS & CDS
- Readiness probe ensures dynamic discovery is complete
- Liveness probe ensures container is running

Data Plane Journey

- Migration from H/W LB to envoy
 - Identified and implemented feature gaps between `envoy` & H/W LB (& contributed upstream)
 - Automation for migrating H/W LB policies to envoy config.
 - Policy Validation - Auto generation of test requests & RUM
- Integration with Propriety Cert Management system
- Data Plane customizations
 - Connection prefetching
 - Drop or RST requests to certain routes based on need
 - Default Upstream Cluster for graceful error handling
- BBR algorithm
- SSL & TCP optimization

Data Plane Journey - Cont. - TCP optimization



Caching on the Edge with Envoy

- Created a custom HTTP filter to support dynamic object caching (Yes. You can!!)
 - Based on URL path, cookies, and request headers
 - Unique for each website user
- Using external cache stores
- Used Envoy's RingHash LB to shard cached data on multiple cache stores
- Better TTFB by order of magnitude (10x)
 - Provides enhanced User Experience
 - Engages User quickly & increases conversion

Future Roadmap

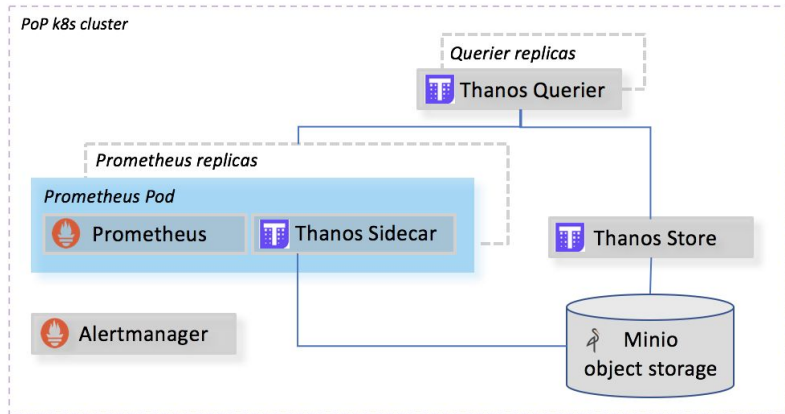
- Enhanced prefetch filter to facilitate Just-In-Time caching

DC vs Edge Proxy vs Caching



Observability

- Prometheus in each PoP
- Central Grafana for visualization
- Envoy specifics
 - Statsd exporter for histogram data
 - FELK for logging data
- Alerting
 - Alertmanager and external checks
 - Static and Dynamic threshold (Anomaly Detection)



Helm based deploy management

- Repeatable deployment across all PoPs
- Helm Charts for the entire PoP stack
- Single source of truth of everything

```
[redacted] (❄️ redacted:14)# helm ls
```

NAME	REVISION	UPDATED	STATUS	CHART	NAMESPACE
alertmanager-prod		:19:40 2018	DEPLOYED	alertmanager-0.1.0	
ats-exporter-prod		:13:38 2018	DEPLOYED	trafficserver-exporter-0.1.0	
elasticsearch-prod		:07:56 2018	DEPLOYED	elasticsearch-0.1.3	
ingress-controller-prod		:56:51 2018	DEPLOYED	ingress-controller-0.1.1	
ksm-prod		:30:55 2018	DEPLOYED	kube-state-metrics-0.1.0	
l4-prod		:02:13 2018	DEPLOYED	l4-0.1.6	
l7- prod		:07:24 2018	DEPLOYED	l7-0.1.14	
l7-prod		:02:06 2018	DEPLOYED	l7-0.1.14	
logstash-prod		:27:20 2018	DEPLOYED	logstash-0.1.7	
minio-prod		:30:13 2018	DEPLOYED	minio-1.6.0	
prometheus-prod		:17:47 2018	DEPLOYED	prometheus-0.2.2	

Learnings and Challenges

Learnings

- Testing is the key
 - HW / SW cross validation
 - RUM to ensure no perf degradation
 - DC and PoP failover test
 - Chaos test
- For smooth workload migration
 - Traffic shift in controlled manner
 - Integration with Legacy systems
 - Cross-function planning and verification

Challenges

- Multi PoP cluster operation
 - Scaling ops with increasing PoPs
 - Drift detection across PoPs
 - Tooling support for cluster management
- Programmable Edge
 - Build Vs. Integrate
 - Howto: act on data
 - Howto: debug and troubleshoot

Thank You !

We are hiring !!!!

bmadhavan@ebay.com

qiuyu@ebay.com