

Building and operating service mesh at mid-size company

Taiki Ono, Cookpad Inc

Agenda

- Background
- Problems
- Introducing and operations
- Key results
- Next challenges

Backgroud

Our business

- "Make everyday cooking fun!"
- Originally started in Japan in 1997
- Operate in over 23 languages, 68 countries
- World largest recipe sharing site:
cookpad.com



Our scale and organization structure

- 90M monthly average user
- ~200 product developers
- ~100 production services
- 3 platform team members
 - 1 for service mesh dev

Each product team owns their products but all of operations are still owned by central SRE team




SRE team



Product team

Technology stack

- Ruby on Rails for both web frontend and backend apps
- Python for ML apps
- Go, Rust for backend apps
- Java for new apps
- Other languages for internal apps



Problems

Operational problems

- Decrease in system reliability
- Hard to troubleshoot and debug distributed services
 - Increase of time detect root causes of incidents
 - Capacity planing

Library approach solutions

- github.com/cookpad/expeditor
 - Ruby library inspired by Netflix's Hystrix
 - Parallel executions, timeouts, retries, circuit breakers
- github.com/cookpad/aws-xray
 - Ruby library for distributed tracing using AWS's X-Ray service

GoPythonJava apps?

- Limitation of library model approach
 - Save resources for product development
- Controlling library versions is hard in a large organization
- Planning to develop our proxy and mixed with consul-template

"Lyft's Envoy: Experiences Operating a Large Service Mesh"

at SRECON America 2017 (March)

Introducing our service mesh

Timeline

- Early 2017: making plan
- Late 2017: building MVP
- Early 2018: generally available

In-house control-plane

- Early 2017: no Istio
- We are using Amazon ECS
- Not to use full features of Envoy
 - Resiliency and observability parts only
- Small start with in-house control-plane, but planned to migrate to future managed services.

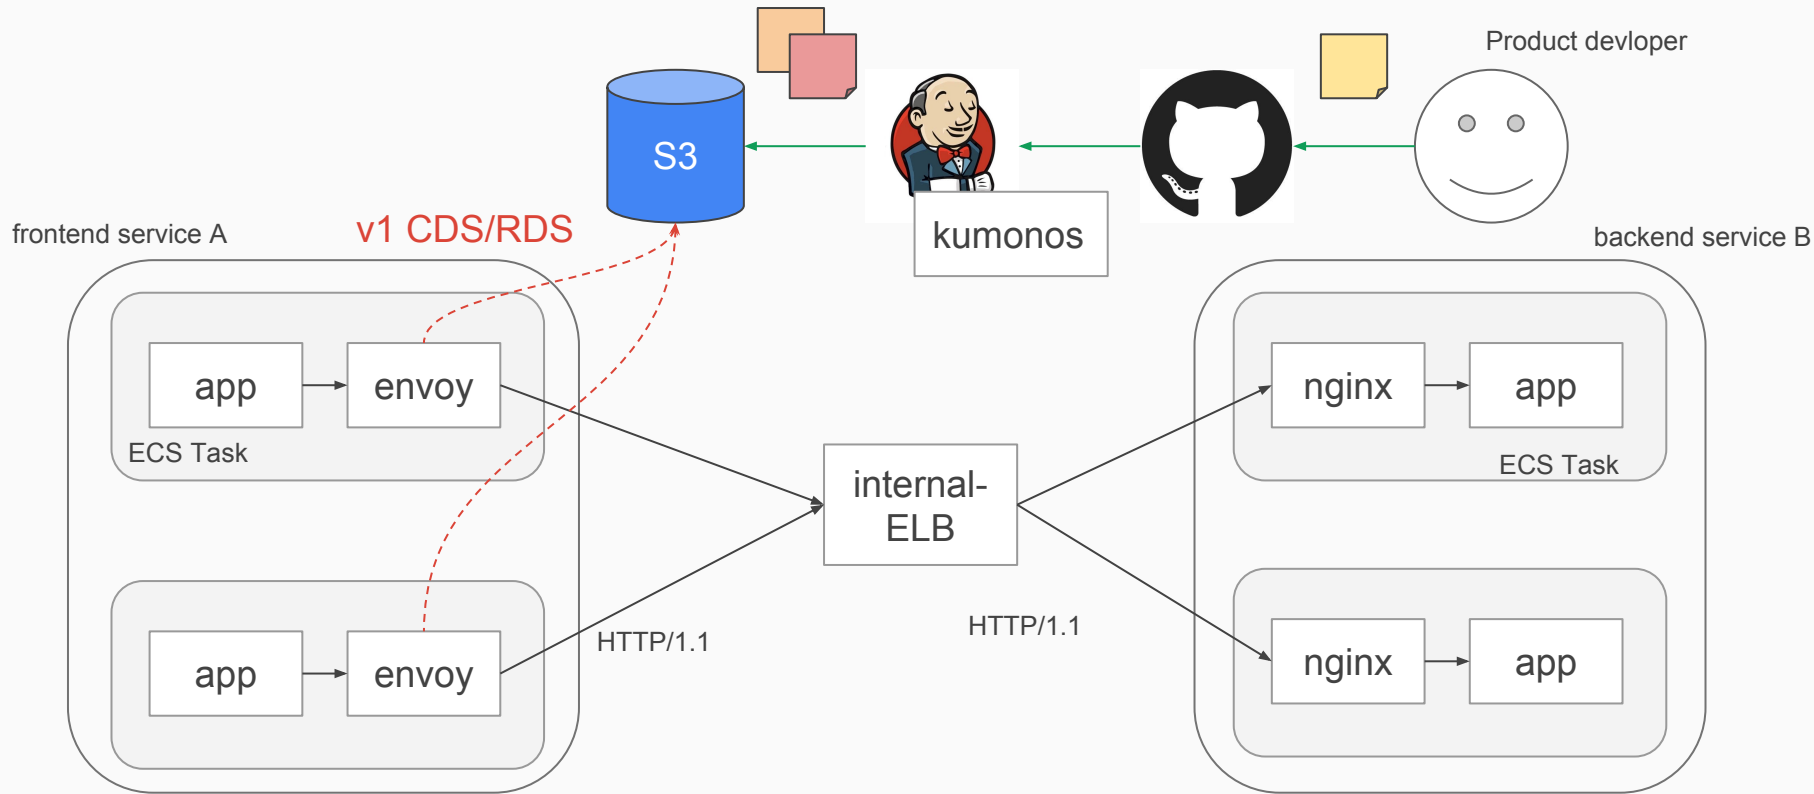
Considerations

- Everyone can view and manage resiliency settings
 - Centrally managed
 - GitOps with code reviews
- All metrics should go into Prometheus
- Low operation cost
 - Less components, use of managed services

Our service mesh components

- **kumonos** (github.com/cookpad/kumonos)
 - v1 xDS response generator
- **sds** (github.com/cookpad/sds)
 - Fork of github.com/lyft/discovery to allow multiple service instances on the same IP address
 - Implements v2 EDS API
- **itacho** (github.com/cookpad/itacho)
 - v2 xDS response generator (CLI tool)
 - v2 xDS REST HTTP POST-GET translation server
 - GitHub#4526 “REST xDS API with HTTP GET requests”

v1 ELB based with v1 xDS



Configuration file

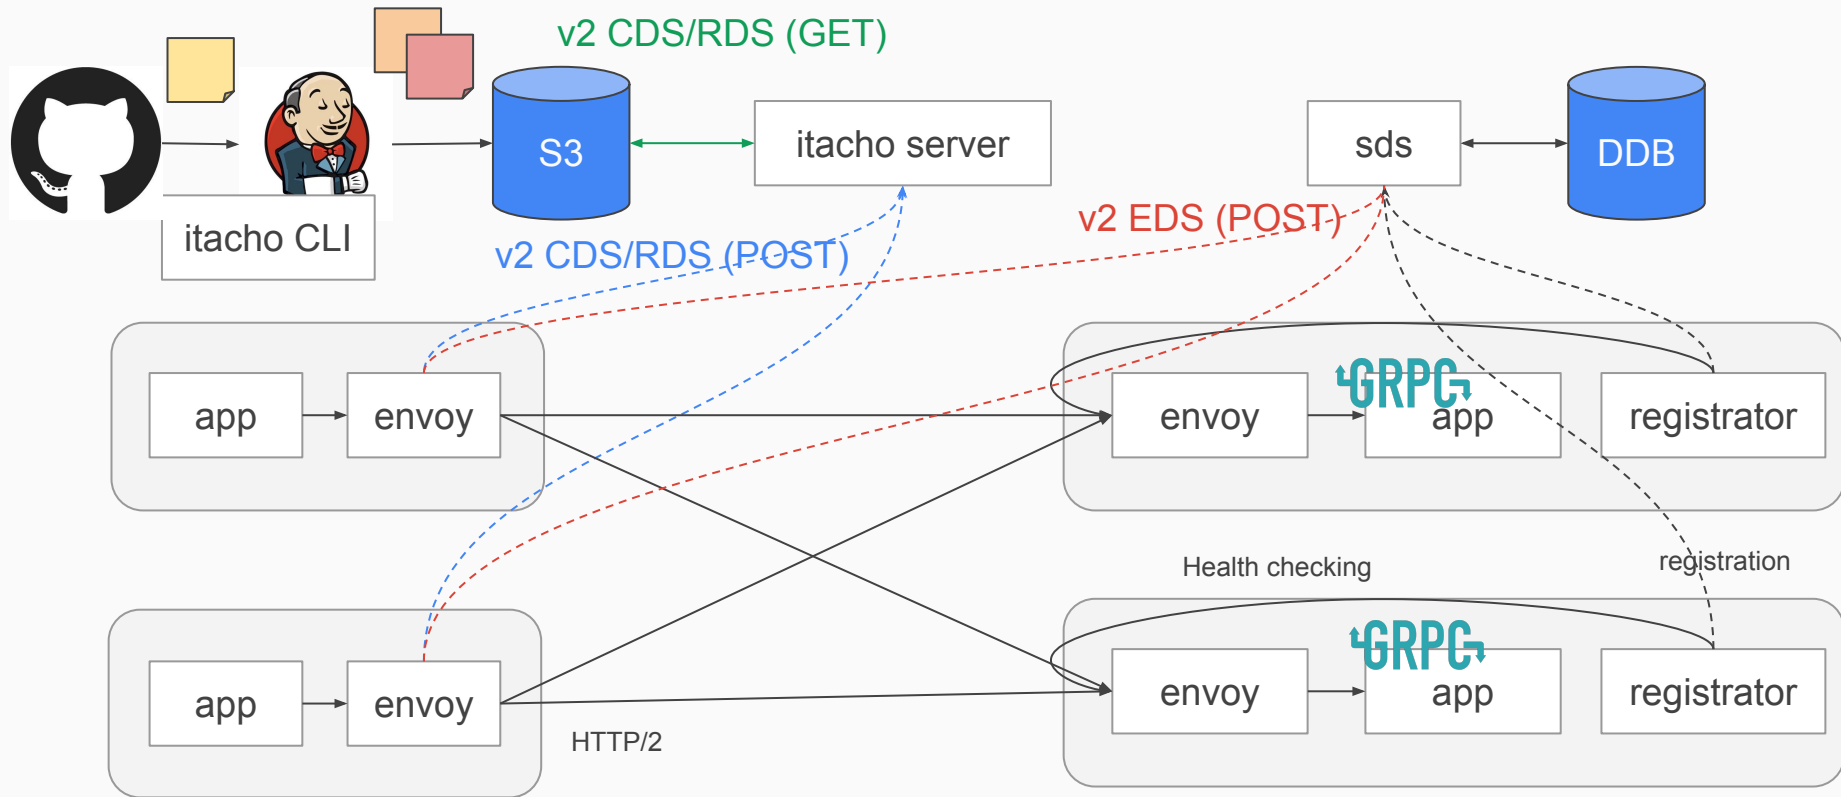
- Single Jsonnet file represents single service configuration
 - 1 service - N upstream dependencies
- Route config
 - Retry, timeouts for paths, domains
 - Auto retry with GET,HEAD routes
- Cluster config
 - DNS name of internal ELB
 - Port, TLS, connect timeout
 - Circuit breaker settings

```
local circuit_breaker = import 'circuit_breaker.libsonnet';  
local routes = import 'routes.libsonnet';
```

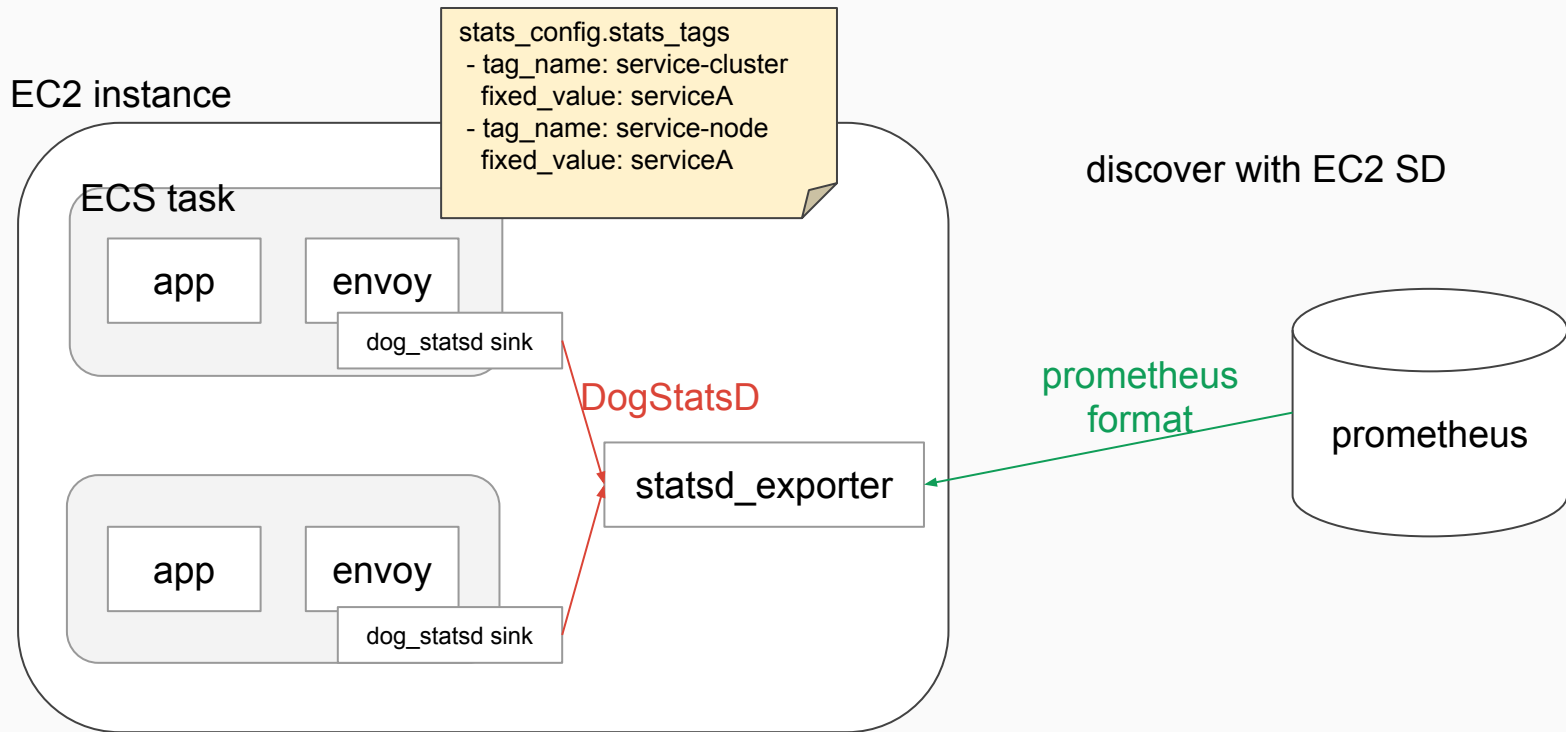
```
{  
  version: 1,  
  dependencies: [  
    {  
      name: "user",  
      cluster_name: "user-development",  
      lb: "user-service.example.com:80",  
      tls: false,  
      connect_timeout_ms: 250,  
      circuit_breaker: circuit_breaker.default,  
      routes: [routes.default],  
    },  
  ],  
}
```



v2 with v2 xDS



Sending metrics



Operation side

Dashboards


- Grafana
 - Per service (1 downstream - N upstreams)
 - Per service-to-service (1 downstream - 1 upstream)
 - Envoy instances
- Netflix's Vizceral
 - github.com/ngghialv/promviz
 - github.com/mjhd-devlion/promviz-front



service

available


Egress status codes by upstream cluster



Egress downstream status codes




Egress RPS by upstream cluster



Egress resp. time by upstream cluster




Egress retries and timeouts by upstream cluster



circuit breaker summary



Egress retry and circuit breaker summary



retry success count

Max:

Avg:

retry overflow counter

Max:

Avg:



Envoy instances

1157

RPS for S3

420

RPS for sds2

202

RPS for sds1

0


CDS update failures

7


RDS update failures

10


CDS config updates



RDS config reloads




CDS update successes




CDS update failures

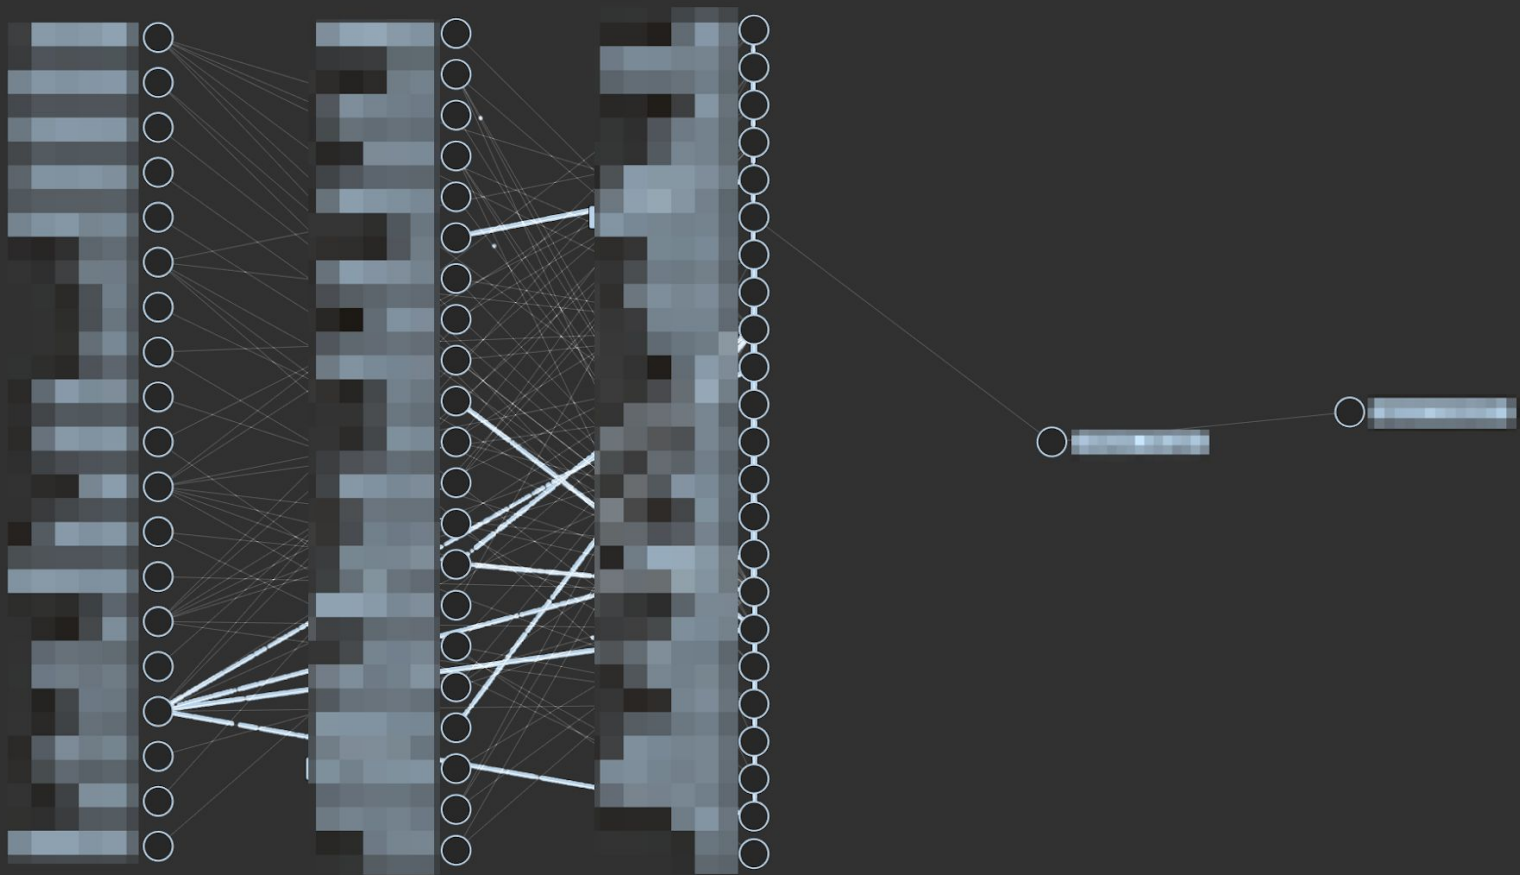



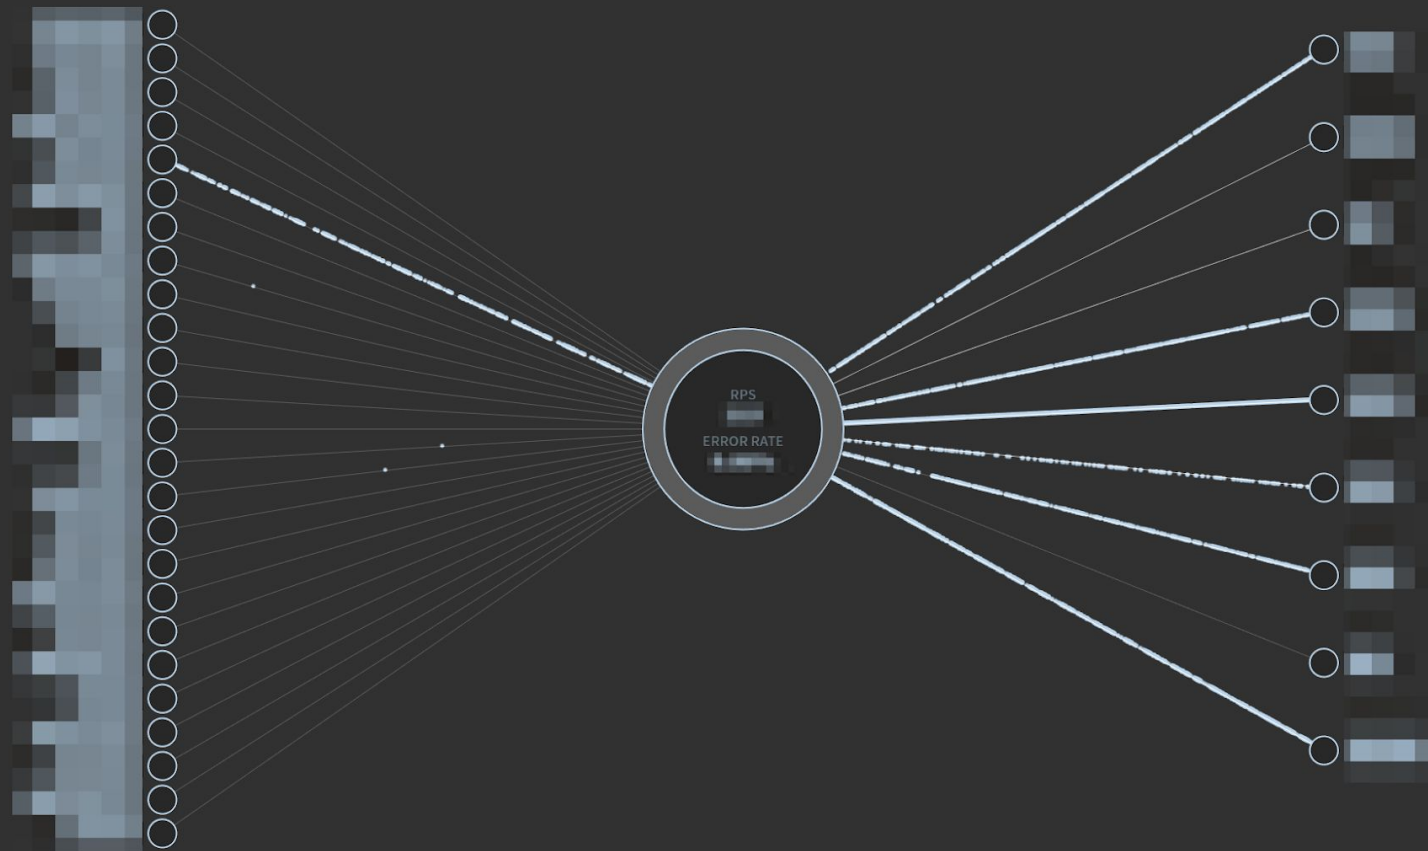
RDS update successes



RDS update failures





 default

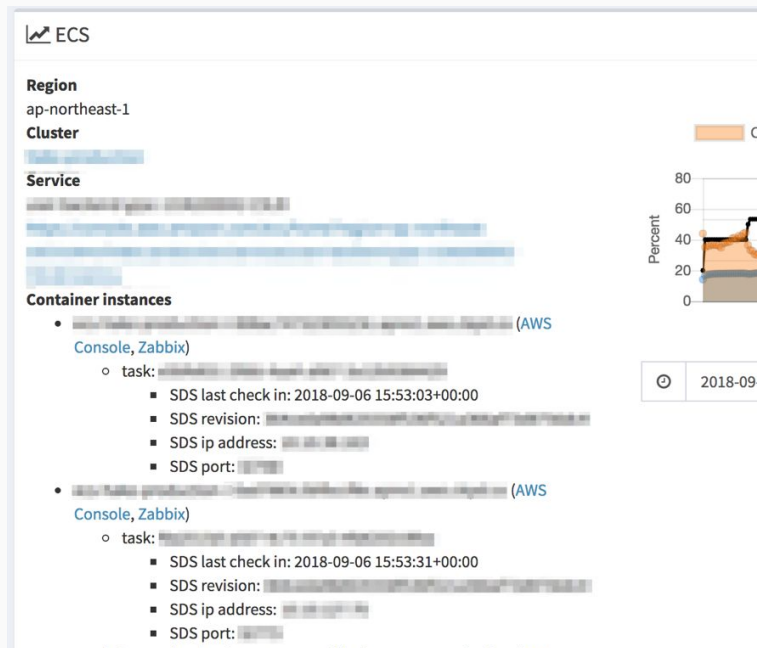
Notices

None

[> INCOMING CONNE](#)[> OUTGOING CONN](#)

Integrate with in-house platform console

- Show service dependencies
- Link to service mesh config file
- Show SDS/EDS registration
- Link to Grafana dashboards
 - Per service dashboard
 - Service-to-service dashboard



Envoy on EC2 instances

- Build and distribute as a in-house deb package
 - Setup instances with configuration management tool like Chef:
github.com/itamae-kitchen/itamae
- Manage Envoy process as a systemd service
- Using hot-restarter.py
 - Generate starter script for each host role

Wait initial xDS fetching

- Sidecar Envoy containers need a few seconds to be up
 - Background jobs are service-in quickly
 - ECS does not have an API to wait the initializing phase
- Wrapper command-line tool
 - github.com/cookpad/wait-side-car
 - Wait until an upstream health check succeed
- Probably move to GitHub#4405

The hard points

- Limitation of ECS and its API
 - Without ELB integration, we need to manage lots of things on deployments.
 - We needed AWS Cloud Map (actually we made almost the same thing in our environment).

Key results

Observability

- Both SRE and product team have confidence in what's happened in service-to-service communication area
 - Visualization of how resiliency mechanism is working
 - Decrease of time to detect root causes around service communication issues
- Necessary to encourage collaboration between multiple product teams

Failure recovery

- Be able to configure proper resiliency setting values with fine-grained metrics
- Eliminates temporal burst of errors from backend services
- Fault isolation: not yet remarkable result

Continuous development of app platform

- Improve application platform without product application deployment
- Increase velocity of platform development team

Next challenges

Next challenges

- **Fault injection platform**
- Distributed tracing
- Auth{z, n}
- More flexibility on traffic control
 - Envoy on edge proxies?
- **Migration to managed services**

Q&A

- Twitter hashtag #EnvoyCon, @taiki45
- Published this slide at: envoyconna18.sched.com